



Miguel de Lemos Dias Rosa Anciães

Bachelor of Computer Science and Engineering

A Trusted and Privacy-Enhanced In-Memory Data Store

Dissertation submitted in partial fulfillment
of the requirements for the degree of

Master of Science in
Computer Science and Engineering

Adviser: Henrique João Lopes Domingos,
DI-FCT-UNL, NOVA LINS

Examination Committee

Chair: Name of the committee chairperson
Rapporteurs: Name of a rapporteur
Name of another rapporteur
Members: Another member of the committee
Yet another member of the committee



FACULDADE DE
CIÊNCIAS E TECNOLOGIA
UNIVERSIDADE NOVA DE LISBOA

November, 2020

A Trusted and Privacy-Enhanced In-Memory Data Store

Copyright © Miguel de Lemos Dias Rosa Anciães, Faculty of Sciences and Technology, NOVA University Lisbon.

The Faculty of Sciences and Technology and the NOVA University Lisbon have the right, perpetual and without geographical boundaries, to file and publish this dissertation through printed copies reproduced on paper or on digital form, or by any other means known or that may be invented, and to disseminate through scientific repositories and admit its copying and distribution for non-commercial, educational or research purposes, as long as credit is given to the author and editor.

Dedictory

ACKNOWLEDGEMENTS

The acknowledgements. You are free to write this section at your own will. However, usually it starts with the institutional acknowledgements (adviser, institution, grants, workmates, ...) and then comes the personal acknowledgements (friends, family, ...).

There is no cloud, it's just someone else's computer

ABSTRACT

The recent advent of hardware-based trusted execution environments provides isolated execution, protected from untrusted operating systems, allowing for the establishment of hardware-shielded trust computing base components. As the processor provides such “shielded” trusted execution environment (TEE), their use will allow users to run applications securely, for example on the remote cloud servers, whose operating systems and hardware are exposed to potentially malicious remote attackers and non-controlled system administrators’ staff. On the other hand, Linux containers managed by Docker or Kubernetes are interesting solutions to provide lower resource footprints, faster and flexible startup times, and higher I/O performance, compared with virtual machines (VM) enabled by hypervisors. However, these solutions suffer from software kernel mechanisms, easier to be compromised in confidentiality and integrity assumptions of supported application data. This dissertation designed, implemented and evaluated a Trusted and Privacy-Enhanced In Memory Data Store, making use of a hardware-shielded containerised OS-library to support its trust-ability assumptions. To support large datasets, requiring data to be mapped outside those hardware-enabled containers, our solution uses partial homomorphic encryption, allowing trusted operations executed in the protected execution environment to manage in-memory always-encrypted data, that can be or not mapped inside the TEE.

Keywords: Hardware Security; Privacy-Enhanced Data Store; Homomorphic Encryption; Isolated Environments; Trusted Computing; Cloud Computing; Virtualisation; Containerisation; Availability; Reliability.

RESUMO

Os recentes avanços de ambientes de execução confiáveis baseados em hardware fornecem execução isolada, protegida contra sistemas operativos não confiáveis, permitindo o estabelecimento de componentes base de computação de confiança protegidos por hardware. Como o processador fornece esses ambientes de execução confiável e "protegida"(TEE), o seu uso permitirá que os utilizadores executem aplicações com segurança, por exemplo em servidores *cloud* remotos, cujos sistemas operativos e hardware estão expostos a ataques potencialmente maliciosos assim como administradores de sistema não controlados. Por outro lado, os *containers* Linux geridos por sistemas *Docker* ou *Kubernetes* são soluções interessantes para poupar recursos físicos, obter tempos de inicialização mais rápidos e flexíveis e maior desempenho de I/O (interfaces de entrada e saída), em comparação com as tradicionais máquinas virtuais (VM) activadas pelos hipervisores. No entanto, essas soluções sofrem com software e mecanismos de kernel mais fáceis de comprometerem os dados das aplicações na sua integridade e privacidade.

Esta dissertação projectou, implementou e avaliou um Sistema de Armazenamento de Dados em Memória Confiável e Focado na Privacidade, utilizando uma biblioteca containerizada e protegida por hardware para suportar as suas suposições de capacidade de confiança. Para oferecer suporte para grandes conjuntos de dados, exigindo assim que os dados sejam mapeados fora dos *containers* seguros pelo hardware, a solução utiliza encriptação homomórfica parcial, permitindo que operações executadas no ambiente de execução protegido façam gestão de dados na memória que estão permanentemente cifrados, estando eles mapeados dentro ou fora dos *containers* seguros.

Palavras-chave: Segurança de Hardware; Armazenamento de Estrutura de Dados em Memória Confiável e Focado na Privacidade; Encriptação Homomórfica, Ambientes Isolados; Computação Segura; Computação em *Cloud*; Virtualização, Containerização; Disponibilidade; Confiabilidade.

CONTENTS

1	Introduction	1
1.1	Context and Motivation	1
1.2	Problem Statement	3
1.3	Objectives and Planned Contributions	3
1.4	Report Organisation	4
2	Related Work	5
2.1	Key-Value Stores	6
2.1.1	Memcached	7
2.1.2	Redis	7
2.1.3	Amazon Dynamo DB	8
2.1.4	Microsoft Azure Cosmos DB	9
2.1.5	Microsoft Azure Cache for Redis	9
2.1.6	Aerospike	9
2.1.7	Discussion	10
2.2	Trusted Computing Environments	11
2.2.1	TPM – Trusted Platform Modules	11
2.2.2	TPM - Enabled Software Attestation	13
2.2.3	HSM – Hardware Security Modules	13
2.2.4	Trusted Execution Environments	14
2.2.5	Intel SGX	14
2.2.6	Sanctum	17
2.2.7	ARM Trust Zone	17
2.2.8	Discussion	18
2.3	TEE/SGX Enabled Key Value Stores	19
2.3.1	Trusted Execution with Intel SGX	19
2.3.2	Circumvention of SGX Limitations	20
2.3.3	SGX-Enabled Secure Databases	21
2.3.4	Discussion	24
2.4	SGX Virtualisation Frameworks	25
2.4.1	KVM-SGX	25
2.4.2	Graphene-SGX	26

CONTENTS

2.4.3	SCONE	26
2.4.4	Asylo	27
2.4.5	Discussion	27
2.5	Related Work Balance and Critical Analysis	28
3	System Model and Design Options	29
3.1	Refinement of Objectives and Contributions	30
3.1.1	SGX Limitations Refinement	30
3.2	Threat Model and Security Properties	30
3.2.1	Adversarial Model Definition	31
3.2.2	System Assumptions	31
3.2.3	Countermeasures for Privacy-Preservation	32
3.3	System Model	33
3.3.1	Key-Value Storage Server	33
3.3.2	Proxy Server	34
3.3.3	Authentication Server	35
3.3.4	Client	35
3.4	System Architecture	35
3.5	Supported Operations	35
3.5.1	Role-Based Authorisation	36
3.5.2	Key-Value Storage Operations	36
3.5.3	Proxy Enabled Operations	37
3.5.4	Attestation	37
3.6	Operation Flow	38
3.7	Summary	39
4	Prototype Implementation	41
4.1	Architecture and Implementation Options	41
4.1.1	Secure Redis	42
4.1.2	Proxy Server	43
4.1.3	Client-based Benchmarks	44
4.1.4	Authentication Server	45
4.1.5	Attestation	45
4.2	Additional Details	47
4.2.1	Protected Memory Check	47
4.2.2	TLS, HTTPS and Certificate Chain	47
4.2.3	Logging and Auditing	47
4.3	Tradeoffs on the Implementation Options	47
4.4	Summary	48
5	Validation and Experimental Evaluation	51
5.1	Testbench Environment	51

5.1.1	Cloud-Based SGX-enabled REDIS Single Instance	51
5.1.2	Cloud-Based SGX-Enabled REDIS Cluster	51
5.2	Benchmarks	51
5.3	Performance evaluation for Testbench 1	51
5.3.1	Latency	51
5.3.2	Throughput	51
5.4	Performance evaluation for Testbench 2	51
5.4.1	Latency	51
5.4.2	Throughput	51
5.5	Evaluation of the Attestation Protocol	51
5.6	Complementary Measurements	51
5.7	Summary and Findings	52
6	Conclusions	53
6.1	Main conclusions	53
6.2	Open Remarks	53
6.3	Future Work Directions	53
	Bibliography	55
	Annexes	63
I	Technologies and Versions	63

LIST OF FIGURES

2.1	Azure Environment Integration	10
2.2	TPM insides	12
2.3	Remote Attestation Procedure	13
2.4	SGX Memory Architecture [45]	15
2.5	SGX Access Control [45]	16
2.6	Arm TrustZone Stack [13]	18
2.7	Server-side components of EnclaveDB	22
2.8	Overview of ShieldStore	24
3.1	System Model Overview	33
3.2	System Architecture Stack	36
3.3	Attestation Model	38
3.4	Operation Flow	39
4.1	Attestation Flow	46

LISTINGS

2.1	Redis Set & Get	6
2.2	How Fast is Redis	7
4.1	Machine Specifications	41

ACRONYMS

ACL	Access Control List
AES	Advanced Encryption Standard
AIK	Attestation Identity Key
API	Application Programming Interface
AWS	Amazon Web Services
CA	Certification Authority
CAS	Configuration and Attestation Service
DBMS	Database Management System
DDoS	Distributed Denial of Service
DoS	Denial of Service
Ecall	Enclave call
ECB	Electronic Code Book
EK	Endorsement Key
EPC	Enclave Page Cache
EPCM	Enclave Page Cache Mapping
EPID	Enhanced Privacy ID
GB	Gigabyte
HSM	Hardware Security Module
HTTP	Hypertext Transfer Protocol
HTTPS	Hypertext Transfer Protocol Secure
I/O	Input/Output
IaaS	Infrastructure as a Service
IoT	Internet of Things

ACRONYMS

IP	Internet Protocol
JVM	Java Virtual Machine
JWT	Json Web Token
KVM	Kernel Virtual Machine
KVS	Key-Value Store
LAS	Local Attestation Service
LRU	Least Recently Used
LSM	Log-Structured Merge Tree
MB	Megabyte
MIM	Man-in-the-middle
Ocall	Out call
OS	Operating System
OTP	One-Time Password
P2P	Peer to Peer
PCR	Platform Configuration Register
PRM	Processor's Reserved Memory
RAM	Random Access Memory
REST	Representational State Transfer
RSA	Rivest–Shamir–Adleman
SaaS	Software as a Service
SASL	Simple Authentication and Security Layer
SDK	Software Development Kit
SGX	Software Guard Extensions
SQL	Structured Query Language
SSL	Secure Sockets Layer
syscall	System call
TB	Terabyte
TCB	Trusted Computing Base

TCE Trusted Computing Environments
TCG Trusted Computing Group
TEE Trusted execution environment
TLS Transport Layer Security
TPM Trusted Platform Module

*

USB Universal Serial Bus

VM Virtual Machine

INTRODUCTION

In this chapter it's presented the context and motivation for this thesis, the main problem statement followed by the goals and objectives and all the planned contributions. In the end, it is presented the structure used in the following chapters of the document.

1.1 Context and Motivation

Cloud computing has gone through many steps, that include grid and utility computing, application service provision and software as a service before reaching the level we know these days. The concept of delivering continuous resources through a global network is rooted in the 1960's. Some experts credit the professor and computer scientist John McCarthy [75] with proposing the concept of computation being delivered as a public utility.

Then, around 1970's the concept of the virtual machine (VM) started to gain popularity as it permitted multiple distinct computing environments to reside on one physical machine.

One of the first major cloud computing moments was the arrival of *salesforce.com* that pioneered the concept of delivering enterprise applications via a simple website. Later, around the 2000's, current big names like Oracle, SAP, Google, Amazon and Microsoft joined the trend and made the cloud world as it is today [2, 3].

Over the past decade, cloud computing has evolve from something service providers told companies they should adopt, to becoming the technology heart of not only major companies, but medium sized enterprises, small start-ups, personal projects and pretty much anyone who works in the computer science world.

Recent studies are foreseeing that 80% of enterprise IT will move to the cloud by 2025 [1]. The array of services provided now are endless and the costs are attractive to

businesses. These services allow developers to only pay for resource usage, and to take advantage of all the power of very large companies. Scalability at request, reliability with daily backups and seamless integration with a lot of other services are some advantages of moving to the cloud. And all of these functionalities without having to manage big infrastructures and a lot of servers, networks, disks, etc... [93].

All of this data and processing happening in someone else's machine started to raise privacy and security concerns. It has become a very attractive target for malicious hackers to attack cloud providers due to the amount of data they process and hold on their services. The best security researchers are always working with the providers to try and mitigate all bugs and vulnerabilities on their very large platforms which has become also a big attack vector. It has been reported by Microsoft, that *"There was a 300 percent increase in Microsoft cloud-based user accounts attacked year-over-year (Q1-2016 to Q1-2017)."* and also *"The number of account sign-ins attempted from malicious IP addresses has increased by 44 percent year over year in Q1-2017."* [70]. Another example published on the Washington Post describes a sophisticated Man-in-the-Middle (MIM) cyber-attack that has targeted Apple's iCloud service in China, in an apparent attempt to collect user names, passwords and other private information [10]. Also, Amazon Web Services has been in 2019 hit by a massive DDoS (Distributed Denial of Service) attack that kept the system down for about 8 hours straight, which can mean thousands of dollars lost by clients [20].

The use of virtual machines to lodge different computing environments on the same physical machine can also raise problems, as explained by the publication *"Seriously, get off my cloud! (...)"* where the researches were able to exploit and obtain RSA (encryption) keys from other VMs deployed in the same physical machine of Amazon EC2 service. This work affirms the need for stronger isolation techniques in public clouds [48].

Docker containers and Kubernetes clusters, used instead and/or alongside traditional VMs, are two of the most popular technologies among cloud providers and cloud server environments these days, and if not managed correctly can become attack vectors. They are already being exploited, most known by the report of Tesla Motors [92], suffering a breach because of an exposed Kubernetes instance [28, 29].

The well known Cambridge Analytica scandal [94] gave the world another perspective about the security guaranteed by the cloud providers, social media and every platform that keeps user's data in a non secure manner. It shows how it can be exploited, sold and manipulated without the owner's consent.

As for Redis as a storage solution, and explained in depth in section 2.1.2, *"Redis should not be publicly exposed as it has no default authentication and all the data is stored in clear text"* and so, according to some studies, there are around seventy two thousand Redis servers available online today, and over 75% of them were compromised and infected with some kind of malware [15, 30, 72].

It's certain that as the cloud environment grows, the motivation for malicious hackers to attack and try to steal information will grow with it, and with all these security breaches increasing year after year, a previously mentioned study also reflects that *"66% of IT*

professionals say security is their greatest concern in adopting an enterprise cloud computing strategy" [1].

1.2 Problem Statement

The problem behind the goals and objectives of this dissertation can be summarised in designing a system, answering the following questions:

Is it feasible to implement a solution for a remote trusted and privacy-enhanced cloud-based key value store system providing strong security and privacy features in a trustworthy solution? Can we design and implement those features with a good trade-off with operational and performance criteria under scalability and reliability guarantees? Can we address the trust-ability by minimising the trust-computing base using protected components running as isolated trusted bases in hardware-shielded trust execution environments? Is it possible to remove the threat of administrators of a cloud platform breaking data privacy? Can we combine the security guarantees with privacy-enhanced in-memory operations supporting big data sets and grained data-structures?

1.3 Objectives and Planned Contributions

The main goal of this dissertation is to design implement and evaluate a privacy enhanced in-memory key-value store to be used as a trusted cloud service with hardware based security features running in a trusted execution environment supported by Intel's SGX technology. To overcome the protected memory limitations and coarse-grained paging as basically supported in SGX memory-management facilities, our solution must support flexible big data sets in a hybrid approach using SGX-mapped protected memory and unprotected memory. Data in unprotected memory will be encrypted and operated in the encrypted form, using partial homomorphic encryption techniques. For implementation purposes we will address our solution to be leveraged from the REDIS technology, enhanced in a architecture using isolated and containerised services running in isolated Intel-SGX trusted-execution-enclaves but supporting all the variations of REDIS-based architectural deployments, e.g.: as a single REDIS server instance, or using replicated instances (with master-slave and master-slave tree-chains, as well as, clustered instances). We will analyse and compare the designed solution in terms of the introduced security benefits and measuring the overheads introduced by the additional security, privacy and trust-ability guarantees.

In this thesis we plan to achieve the following contributions:

- **Design and implementation of the cloud-enabled privacy-enhanced solution**, with all-in-the-box planned features as described above, and able to be used as a "cloud-platform as a service" solution, providing:

- **Trust-ability, security and privacy** properties with the following guarantees: high availability, built-in replication, LRU eviction model, in-memory operations on encrypted big data sets and complementary options for protected on-disk persistence.
- **Software attestation** guarantees provided to clients (users) in order to validate the correctness and integrity of the remote software stack providing the solution.
- **Multiple Replication Mechanisms** based on the same secure solution to analyse how these types of replication (centralised solution, a Master-Slave architecture and a clustering solution) will be impacted by the additional security features.
- **Drastically reduce TCB** in the remote cloud provider by removing the millions and millions of lines of code implementing the hypervisors and operating systems used in their infrastructures thus creating a **truly isolated system** by leveraging Intel's [SGX](#) technology to create a shielded and trusted execution environment in a remote cloud provider.
- **Complete analysis report** of the different solutions of replication and security levels, comparing a normal non-secure solution with the the privacy-enhanced implementation along with evaluation of overheads and trade-offs introduced by the additional security mechanisms.

1.4 Report Organisation

The remaining of the report is organised as follows:

Chapter two presents the topic background, related work and initial research performed for this thesis, including relevant contributions and similar solutions existing in current days.

Chapter three discusses the approach to the elaboration phase by describing the implemented system architecture and technologies that were used. It provides a in-depth explanation of how the goals and objectives of this dissertation were achieved.

Chapter four provides a planned timeline to be followed throughout the elaboration of this thesis including a breakdown of the work-plan by weeks and categories from the beginning to the thesis delivery and presentation.

CHAPTER 2

RELATED WORK

This chapter presents and briefly discusses the related work and the study performed beforehand in order to guide and give some context to the reader. It will present work that was used as the basis of this thesis, existent technologies and their relation with this project, and some comparisons between those existing technologies, the problem addressed in this thesis and the solutions proposed to solve, or better address, those very same problems.

First, in section 2.1 we explain and discuss for the first time the definition of a Key-Value Store. We present some use cases, current technology available, their differences and most importantly their security models and concerns. Having discussed the software, section 2.2 will then address the environment on where that software will run - hardware. It explains and present the different ways to secure and authenticate the hardware, prevent hardware-based attacks and discuss some of the current products available and how they will be used across this thesis. Section 2.3 will then make the bridge between software and hardware. It explains how Key-Value stores are currently being run on secure environments. It discusses how software and hardware work together to achieve a secure environment for an application to run. Section 2.4 will discuss how virtualisation can allow for faster deployments and easier implementation of secure code for the complicated frameworks of secure hardware such as Intel's SGX processors. To conclude the chapter, section 2.5 will combine the information of every sub-chapter and analyse it with a bigger perspective and better knowledge of the theme.

Section 2.3 is considered to be the **main core** investigation and directly related to the work planned for this thesis. As for the other sections, they provide a background knowledge necessary for understanding of the core of this dissertation.

Along the next chapter we summarise the main relevant ideas that can be retained from each section for our objectives and expected goals.

2.1 Key-Value Stores

Key value stores are the simplest form of what computer scientists call a database. The simplicity lies on associating a value to a certain key and storing that pair, as well as retrieving the values of known keys. [53]

Listing 2.1: Redis Set & Get

```
1 redis> SET mykey "Hello"
2 "OK"
3 redis> GET mykey
4 "Hello"
```

Is this simplicity that makes this technology very attractive to developers. The ease of use, its high performance and speed are key aspects in favour of this technologies. However, simply working with keys and values might not be enough to more complex applications, and that is why Key-Value store product developers are introducing new features in order to make them appealing to a broader mass of users, always keeping them lightweight and fast.

For that lightweight and fast attributes, most of the key-value stores work in the computer memory. This allows fast read and write operations as opposed to persistent disk storage. Although they work mainly in memory, most of the solutions offer some persistent mechanism so we can make use of its performance but still persist data in case of a disaster, server failure or any crash.

KVSs have been evolving for years and some are now more than a single key-value store module. A lot of them are now supporting a multi-model storage. Meaning that a value can be more than a single integer or a string. For example, Redis [76] as a multi-model store is not only a key-value store, but also [78]:

- **Document Store** - *"nonrelational database that is designed to store and query data as JSON-like documents"* [35]
- **Graph DBMS** - *"Graph databases are purpose-built to store and navigate relationships. Use nodes to store data entities, and edges to store relationships between entities"* [41]
- **Search Engine** - *"nonrelational database that is dedicated to the search of data content. Use indexes to categorize the similar characteristics among data"* [83]
- **Time Series DBMS** - *"Provides optimum support for working with time-dependent data. Each entry has a timestamp, the data arrives in time order and time represents a primary axis for the information."* [95]

So, the **KVS** world is becoming more and more versatile as the years pass.

In the next subsections its discussed and presented the overview of the current **KVS** technology. We picked the some top KVSs technologies nowadays according to db-engines [54] website.

2.1.1 Memcached

Memcached [65] is a free and open source key-value store released in 2003. It is described as a high performance distributed memory object caching system.

It is design to hold small chunks of data (strings and objects) to work as a cache for results of database calls, API calls, or page rendering. Its biggest use case is for use in speeding up dynamic web applications by alleviating database load.

This system lies on the simpler key-value store spectrum. It takes advantages of the simplicity of a key-value store to edge ease of development, and solving many problems facing large data caches. Its API is available for most popular languages. It has a [LRU](#) eviction technique which means that items will expire after a specified amount of time if not used.

When it comes to system replication, availability and reliability, Memcached has an interesting approach. In order to keep it blazing fast, there is no communication between server instances in a cluster. Memcached servers are unaware of each other. There is no crosstalk, no synchronisation, no broadcasting, no replication. Adding servers will only increase the available memory.

As for its security context, Memcached spends very little, if any, effort in securing the systems for random internet connections. The servers only have support for [SASL](#) [81] authentication mechanism. This method of authentication is not implemented as end-to-end encryption, it only provides restriction access to the daemon, but it does not hide communications over the network. That means it is not meant to be exposed to the internet or to any untrusted users [66].

2.1.2 Redis

Redis [76] is an in-memory data structure store that can be used as a database, cache and also a message broker. Redis focuses on performance, so most of its decisions prioritise high performance and very low latency.

It has been benchmarked as the world's fastest database [79] and together with a their multi-model and its rich set of operations that can be performed over data it has been the leading key-value store according to use and popularity for a multiple set of years [54].

Listing 2.2: How Fast is Redis

```
1 $ redis-benchmark -t set -r 100000 -n 1000000
2 ===== SET =====
3 1000000 requests completed in 8.78 seconds
4 50 parallel clients
5 3 bytes payload
6 keep alive: 1
7
8 99.59% <= 1 milliseconds
9 99.98% <= 2 milliseconds
10 100.00% <= 2 milliseconds
```

11	113934.14 requests per second
----	-------------------------------

As said before, Redis is now not a simple [KVS](#). It supports data structures such as strings, hashes, lists, sets, sorted sets with range queries, bitmaps, hyperloglogs, geospatial indexes with radius queries and streams. It also has built-in replication, server side scripting, [LRU](#) eviction, concept of transactions and different levels of persistence. It provides high availability and automatic partitioning as well.

Redis provides two modes of replication/availability. The master slave form of replication works with a single node (master node) where all writes occurring will be replicated to the other Redis instances (slave nodes). Writes on nodes other than the master will not be replicated. Redis provides a read-only setting that can be applied to slave nodes to prevent states differences between instances. The cluster mode of replication consists on partitioning the data between multiple master nodes. This mechanism does not provide replication as each data object is stored on one master node only, but will allow for lower dataset sizes on each node and therefore faster response times. This method is usually mixed with the first one to provide both performance and replication features, by adding slave node to each of the masters.

Security is not Redis' primary concern (just like others). *"In general, Redis is not optimised for maximum security but for maximum performance and simplicity"* [80]. It is design to be accessed by trusted clients inside trusted networks. This means that it is not supposed to be publicly exposed. Redis (in its latest release 6.0.6 [77]) now implements an [ACL](#) policy that allow the configuration of multiple users with different permissions over the dataset. It also added support for [SSL](#) network communication security.

There are a few other security concerns that Redis addresses, but has we can now start to see, in this types of stores, security falls behind performance and usability.

2.1.3 Amazon Dynamo DB

Amazon Dynamo DB [6] is a fully managed NoSQL database service. It is a key-value store and a document store that is built based on the dynamo paper [31]. This paper describes a [P2P](#) (peer-to-peer) network with high availability, eventual consistency and very easily scalable. It also successful handles server and data center failures and network partitions.

Amazon builds on this paper and offers DynamoDB as a service in their platform. It is a hosted system in the Amazon Web Services [8] infrastructure and it is fully managed. That means no need for low level server configurations or maintenance. It is all managed by the [AWS](#) team and offered to the user with a nice configuration interface. It also means that it has built-in security, backup and restore and in-memory caching for internet-scale applications. Also, it offers seamless scalability by increasing the number of nodes/servers according to current traffic received by the application on any given time.

This technology focuses more on high availability but also achieves very high performances and very low latency and being fully managed it also takes advantages of the [AWS](#)

infrastructure full power. It currently sits second on the db-engines [54] most popular ranking.

2.1.4 Microsoft Azure Cosmos DB

Microsoft Azure Cosmos DB [68] is a fully managed database service provided by Microsoft Azure [69]. This service provides a global distributed, horizontally scalable, multi-model database. Its multi-model architecture can work as a key-value store, a Document Store, a graph DBMS and a wide column store.

It's very proud and excels in the ease of global scale with the system call *Turnkey global distribution*, providing transparent multi-master replication and a set of user configurable consistency options. It also strongly advertises a *Multi-Model Multi-API* feature where you can use multiple data types on this single database service. Cosmos DB automatically indexes all data and allows the user to use various NoSQL APIs to query the data.

As a fully managed service, Cosmos DB makes use, in the background, of the large infrastructure with almost unlimited resources and capabilities provided by Microsoft, which means it also has built-in security, fail-over mechanisms for disaster recovery, and high performance with single digit read and write latencies.

2.1.5 Microsoft Azure Cache for Redis

Microsoft Azure Cache for Redis [67] is a service provided by Microsoft Azure that joins the open source world of Redis with the commercial side of a fully managed and hosted platform.

It uses at its core the Redis server technology and provides ease of deployment and management, built-in global replication, Azures' infrastructure security and flexible scaling and Redis superior throughput and low latency performance.

Being in the Azure ecosystem provides nice integration with all Azures' services as shown in figure 2.1.

2.1.6 Aerospike

Aerospike [4] is an enterprise-grade, high performance Key-Value Store. It is another KVS technology currently available today. It promises a philosophy of "*no data loss*" through Strong Consistency. Normal systems trade requiring this type of consistency usually trade performance for data integrity but Aerospike allows it with minimal performance loss. That means it can be used for example in banking payments, retail and telecommunications use cases.

It also provides a dynamic cluster management and unique flexible storage. That enables very easy deployments and particularly very easy scalability, so it is able to meet any data volume needs and still maintaining low latencies across that wide range of data volumes, from low volumes until hundreds TB of data.



Figure 2.1: Azure Environment Integration

As for security, it includes (the enterprise version) a database access management and audit trail logs. It also includes transport level encryption for client-server traffic and cross-datacenter traffic [5].

2.1.7 Discussion

In this chapter when gather information about the overview of the current Key-Value Store. We can conclude that the most important feature of this technology is the performance and all of the above products mentioned do focus on that characteristic. Some of them even compromise in another features to achieve the best performance possible. Security is not the main concern and the most used measures in the current technologies being securities implementations at the network and transport level by using [TLS](#) and also full disk encryption.

Network and transport layer security is a must when implementing any system, and this thesis will also use those standards.

As for full disk encryption on the server, it opens up some attack vectors. Full disk encryption means that random users will not be able to query the data but credentialed users can. Although, anyone with full access to the database, for example database operators or/and administrators, can decrypt and access all information. This creates a risk of privacy breaking due to hackers wielding stolen credentials, rogue insiders who have been granted more access than they need or the well known honest-but-curious adversary model, where an administrator with full credentials does not have bad intentions, but, driven by curiosity, access information therefore breaking data privacy. A cloud based

[KVS](#) service like the ones talked above, this type of vulnerabilities can be a major concern for a use case with very sensitive data since the server would be off premises, there is no control over it when it comes to privacy of data.

This thesis will implemented a system based on Redis, the most popular and used Key-Value Store currently used and will try to solve some of the problems with security described above. It will compare the principle feature of a [KVS](#), the performance, of a simple and normal Redis server and a privacy-enhanced Redis solution so the user can calculate the trade-off between performance and security and applied the correspondent solution to their own use case.

2.2 Trusted Computing Environments

Modern data processing services hosted in the cloud are under constant attack from malicious system administrators, server administrators and hackers who exploit bugs on applications, operating systems or even the hypervisor. However, current days shows a massive trend of business moving to the cloud infrastructure looking for easy deployment, managed services with built-in replication and fault tolerance, fast and trivial scaling and predicted costs.

With more and more data exposed in the cloud, hackers have a bigger desire to exploit and look for vulnerabilities. This results in frequent data breaches that reduce trust in online services. The need for cloud providers to ensure a level of security and trust to make the user comfortable of moving its data to the cloud has never been bigger, and with that need some solutions in the form of Trusted Computing Environments ([TCE](#)) appeared.

Trusted Computing is a concept that strives to provide strong confidentiality and integrity guarantees for applications running on untrusted platforms. It forces a certain machine to behave an expected way even if running on a remote or machine that is out of our control.

[TCE](#) will also provide a decrease of the Trusted Computing Base ([TCB](#)) - the amount of components that the application needs to trust in order to run smoothly. By isolating the service running on this trusted environments (limiting the set of instructions available and encrypting data), it prevents the operating system, the hypervisor and even malicious system administrators (three components normally on the [TCB](#)) to break data confidentiality and integrity within this environments.

There are a few hardware/software based solutions to achieved a trusted computing environment, and they will be explained in the next sections.

2.2.1 TPM – Trusted Platform Modules

A Trusted Platform Module, also known as a [TPM](#) is a technology proposed by the Trusted Computing Group ([TCG](#)) designed to provide hardware-based security related functions.

It's a chip embedded into the motherboard and includes multiple security mechanisms to make it tamper resistant to physical harm and malicious software is unable to mess with its security features [96]. Some key advantages of using TPMs are:

- Generate, store, and limit the use of cryptographic keys
- Platform identity by using the TPM's unique RSA key, which is burned into itself also known as Endorsement Key (EK) and never leaves the TPM.
- Help ensure platform integrity by taking and storing security measurements.

Figure 2.2 shows the main components and services provided by a TPM module. As shown in the figure, all of them only have one access point I/O which is protected and safely managed by the TPM execution engine.




Figure 2.2: TPM insides

With all the components described by figure 2.2, TPMs provide their main TPM features: Encryption, Authenticated Boot and Attestation.

The first feature is used for every security and confidentiality aspects, mainly generating cryptographic keys, encrypting, signing and hashing data with secure standard algorithms melted in the module.

Authenticated Boot is the ability to boot the OS in stages, assuring that each portion of OS, as it is loaded, is a version trusted and approved for use, detecting hardware and software changes on every stage to verify if the code loaded can be trusted. This boot sequence happens with the help Platform Configuration Registers (PCR) that store the trusted software hashes.

The attestation feature is a way for a client to remotely check the state of a machine and will be further explained in the next subsection.

2.2.2 TPM - Enabled Software Attestation

The remote attestation feature of a **TPM** is the ability of a program to authenticate itself against external verifiers. Is a mechanism that allows a remote party to verify the internal state of the OS or another software and decided whether or not that piece of software is intact and trustworthy. The verifier can trust that the attestation data is accurate and not tampered with because it is signed by the internal key of the **TPM**, a special key known as the Attestation Identity Key, known from now on as **AIK** [19].

A remote attestation procedure is described in image 2.3 [18]:



Figure 2.3: Remote Attestation Procedure

1. The application "A" generates a public/private key pair PK_A & SK_A and asks the **TPM** to certify it.
2. The TPM computes a hash value $\#A$ of the executable code of program "A".
3. The TPM creates a certification including PK_A and $\#A$ and signs it with the attestation identity key SK_{AIK} .
4. When application "A" wishes to authenticate itself to a remote party, it sends the certificate of its public key and hash value $\#A$ along with a certificate issued to the TPM by a trusted certification authority (CA).
5. The remote party verifies the certificate chain.
6. The remote party looks $\#A$ up in a database which maps hash values to trust levels.
7. If application "A" is deemed trustworthy, we continue the communication, probably by using PK_A to establish a session key.

2.2.3 HSM – Hardware Security Modules

An (**HSM**) hardware security module is normally an external module that can be added to a system, in form of a **USB** device or a component living in a secure network as a

trusted server, instead of being embedded into the motherboard like a [TPM](#). It provides a dedicated system of hardware enable accelerated cryptographic functions like encryption, decryption, key generation and signing capabilities [46]

What makes this devices so secure, like the [TPM](#), is it can't be interfered with by external code, and it provides an array of protective mechanisms to detect and prevent external physical tampering like drill protection foil, resin-embedded chips as well as temperature and voltage sensors. Any detection of tampering will result in an alarm as well as countermeasures by the applications installed inside. [24].

[HSM](#) can have various applications and can be used in simple forms for example a specific bank dongle that generates [OTP](#) (one-time password) for accessing your account or be a big corporation and enterprise appliance in various industries, e-health, automotive and [IoT](#) systems.

2.2.4 Trusted Execution Environments

A Trusted Execution Environment ([TEE](#)) is an abstraction that describes a machine capable of executing a given program P in isolation, i.e. whose output is determined by the initial state of P and a set of defined inputs given into the [TEE](#) (Barbosa et al., 2016).

It is a secure area of the main processor that ensures sensitive data and code loaded inside is stored, processed and protected in an isolated and trusted environment. As such, it offers protection from software attacks even the ones generated in the operating system.

A [TEE](#) guarantees that:

- The code loaded in the environment is authentic and was not tampered by an attacker.
- All system state is correct (CPU registers, memory and sensitive I/O).
- The code, all data generated and runtime state is confidential and stored persistently.

The threat model of a [TEE](#) should include all software attacks and the physical attacks performed on the main memory and its non-volatile memory.

"There are many interpretations of what is meant by Trust. In the [TEE](#) it is used to imply that you may have a higher level of trust in validity, isolation and access control in items (assets) stored in this space, when compared to more general purpose software environments"[100].

2.2.5 Intel SGX

"Intel® Software Guard Extensions (Intel® SGX) is a set of instructions that increases the security of application code and data, giving them more protection from disclosure or modification."[49].

These set of instructions are one of the latest iterations of trusted computing solutions and designs that tries to tackle the problem of securing remote computations by leveraging secure hardware on the remote host machine. The SGX processor enables a secure container called enclave which protects the confidentiality and integrity of the execution, such as code and data while relying on software attestation mechanisms.

A **SGX** can be though as a reverse sandbox. With a sandbox you are trying to protect the system from your application, but with **SGX** you are trying to do the opposite and protect the application from the system. The system can be the **OS**, the hypervisor, the BIOS, the firmware or even the drivers [88].

A **SGX** enabled application is broken into two parts, the untrusted and trusted parts. The trusted part of the application is all the processing that deals with any sensitive data the application is handling. This part will be run inside enclaves and be stored in protected memory. The rest will live in normal memory and not be protected.

It provides this kind of security from the hardware by isolating all the private data from the outside, placing it into a restricted area of the memory called the **PRM** more specifically in the **EPC** (Enclave Page Cache) as shown on figure 2.4. The **PRM** is a zone of the **RAM** with guaranteed access management by the CPU where it will deny every external access and only allow access through the associated enclave. This region of the memory is also known as the private/protected memory or the trusted part of the application. The data managed in **EPCs** are mapped in plaintext, only in on-chip caches. They are encrypted and integrity-protected when they are mapped in the external (not protected) memory.




Figure 2.4: **SGX** Memory Architecture [45]

Because the **SGX** enclaves execute within the virtual address space of a process, the translation of enclave addresses must be trusted. However, since it is the **OS** that manages the translation between physical and virtual addresses (and the **OS** cannot be trusted), **SGX** maintains an internal data structure called the **EPCM** (Enclave Page Cache Mapping)

which tracks the referred mapping as well as the information described on figure 2.4 [71].

With all this information about enclave pages, the processor can now performed a controlled access management to the enclave page cache described in figure 2.5, where it will denied access not only from outside the enclave but as well as from enclaves that do not own the page of memory request creating then, an isolated memory region.




Figure 2.5: SGX Access Control [45]

The **SGX** will create an enclave when sensitive code needs to run by a specific **SGX** CPU instruction (ECREATE) and will create a unique instance of an enclave, establishing the linear address range and load the sensitive code into an **EPC** inside the protected memory. Once all pages are loaded into the EPC, and the loading is complete, an authentication hash is computed and is available for remote attestation so a user can verify that the code running in the enclave has not been tampered with.

An enclave must expose an **API** for the application to call in and advertise what services provided by the untrusted domain are needed. This is the definition of an interface boundary between the untrusted part of the code and the enclave and it is how they communicate. An **Ecall** is a function that the untrusted part application can call to execute some code inside the enclave and since it exposes a sensitive interface, to reduce the enclave attack surface, the number of **Ecalls** should be limited. On the other hand, the **Ocall** is a function that an enclave can call to reach a service/interface outside the enclave, on the untrusted **OS**. Again, calling some service out of the enclave can carry additional security risks and should be as minimal as possible [51].

When running, the execution always happens in protected mode, and to prevent data leaking, the CPU will not directly address an interrupt, fault or VM exit, but will instead emit another specific instruction (EEXIT) to properly exit the enclave, save CPU state into the enclave and only then will service the fault.

With all this properties, Intel® through **SGX** set of instructions and implementation

tries to achieve a secure and trusted environment with guarantees of code and data isolation, confidentiality and integrity from attackers such as the OS, hypervisor, any hardware and even physical attacks [26].

2.2.6 Sanctum

Sanctum [27] is an open-source project that shares the same goal as Intel SGX, providing strong provable software isolation to protect the data from external hardware and software, but claims to be simpler and protect against indirect attacks called side channel attacks [58] such as cache timing attacks [23] that have been known to exist in SGX [40, 82]. These are additional software attacks that can infer private information by analysing a program's memory access patterns.

Following the minimal and simple concepts, it uses minimal invasive hardware and it does not require any modifications to the CPU major blocks, but only adds hardware to the interfaces between blocks. This allows for a respectable overhead by maintaining normal clock speeds as it does not modify the CPU core critical execution path.

Sanctum project builds on the SGX programming model and implements an architecture that deviates as little as possible from the one built by Intel. Although it differs from SGX by implementing the enclaves via a small combination of hardware extensions to RISC-V (an open source set of CPU instructions [99]) and a trusted piece of software called the security monitor, as SGX implements them via hardware microcode and presents a set of CPU instructions to manage the enclaves.

This security monitor is the core of the project and configures the hardware to enforce low-level rules that control the enclaves's access policies. As explained in the Sanctum paper, *"the security monitor checks the system software's allocation decisions for correctness and commits them into the hardware's configuration registers"*. One of the examples and the main points of upgrade compared to SGX is that Sanctum keeps the enclave page tables inside enclave memory, protecting the system against the timing attacks referred above by keeping the page table dirty and accessed bits private. Their hardware extensions make sure that enclave page tables only point to enclave memory and untrusted OS tables only point to OS memory regions and never to enclave private memory.

Sanctum is also open to the public which makes easier for security researchers to audit and find vulnerabilities and to further encourage the analysis of the code, Sanctum security monitor is written in portable C++ code and can be used across different CPU implementations.

2.2.7 ARM Trust Zone

ARM Trust Zone [12] is a technology that offers a system wide approach to security based on hardware enforced isolation built into the CPU [11]. The principle of the technology is to separate the trusted and untrusted by two virtual processors backed by hardware

access control. The two states are referred as worlds, where the first is called the secure world (SW) and the other is the normal world (NW) like figure 2.6 shows.

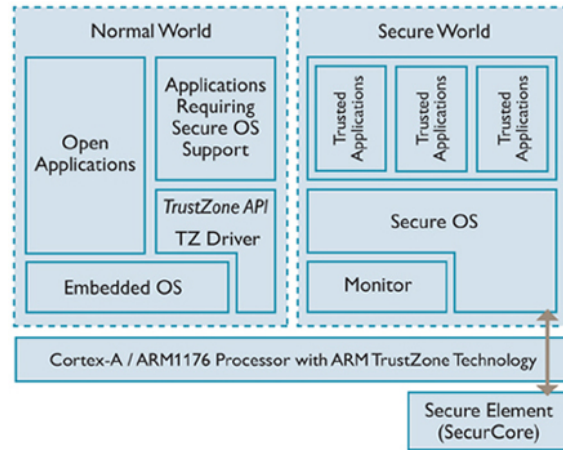


Figure 2.6: Arm TrustZone Stack [13]

The non-secure world (or normal world) is where the OS and most of the software and applications will be running, as for the secure world is where more secure and sensitive software will run and will ensure that vital information is not intercepted by a third party. The security is enforced because each of the worlds acts isolated from the other as a runtime environment with separated resources such as memory, processor, cache, controllers, interrupts. The ARM hardware has separate copies of the registers for each world and cross-world register access is blocked. However the Secure Monitor shown in figure 2.6 can access non secure registers while running in secure world. This means that the monitor can then implement context switching between both worlds.

When in Normal World, the application calls a specific ARM instruction call SMC (Secure Monitor Call) to call back inside the secure world and execute in code in a secure manner.

By keeping the worlds separated from each other, the ARM TrustZone can keep applications running in secure mode isolated from the normal world applications such as the OS and thus achieving another implementation of a TEE.

2.2.8 Discussion

The TCB (Trusted Computing Base) is the set of computer components (hardware), software and data that we need to trust and deem as not malicious in order to use a system. It's a group of various elements that are critical to a system's security in a way that any bug or vulnerability occurring from inside the TCB components might compromise the security and privacy of the entire system. On the other hand, security flaws and bugs from outside of the TCB should not become a security issue.

Hardware security like TPMs and TEE technologies have the goal to drastically reduce the TCB of a system, for example, remove the OS from the trusted base to make sure that

a remote compromised machine vulnerability does not affect the security and privacy of a application or system.

By isolating the program from outside and uncontrollable sources like cloud infrastructures [OSs](#), hypervisors and hardware we can more safely deploy sensitive applications to those cloud providers. That is the goal of all different implementations of trusted execution environments like the Intel [SGX](#) and the ARM TrustZone.

Although not free of some problems, [SGX](#) implementation of a [TEE](#) seems to be the most accepted technology, with serious and skilled developers and security researches always working to mitigate any vulnerability in order to create a truly trusted and isolated environment.

2.3 TEE/SGX Enabled Key Value Stores

There has been an increase trend from developers to move their applications to the cloud. It provides dynamically and almost seamlessly scaling with predict cost. Although it also means that users need to rely on the cloud providers for securing and maintaining the integrity of their applications. That means the user must trust not only the provider's staff but also its globally distributed software and hardware not to expose their private data. Today's cloud providers only aim to protect their privileged code from the untrusted code (the user's code) and do not provide any guarantees about the opposite scenario.

To mitigate this use case, and after studying and discussing the Key-Value stores technologies and also the trusted platform modules as well as the trusted execution environments, in this chapter, it will be presented how are this two topics being combined and used together.

It will be more focused on the Intel SGX platform as it is the one that will be used throughout this thesis. Currently, there are a number of databases who leverage this technology to provide a more secure environment and service. In the this chapter it's presented how they work and operate, discussed the differences between them and also the how the work planned to be performed on this thesis will solve some of the problems and caveats.

2.3.1 Trusted Execution with Intel SGX

As explained before, Intel SGX provides a trusted execution environment by running code inside the enclaves. It creates an isolated environment where we can run some instructions as securely as possible, without [OS](#) intervention.

Key-Value Stores and other database type systems can leverage this secure and isolated environment to perform queries on very sensitive data that would otherwise be vulnerable to some attacks. There are a few techniques currently implemented to use isolated environments. Maintaining an encrypted database and using enclaves cryptographic capabilities to decrypt data and perform queries on plain text with the assurance

of no data leaking is a possible use case. Also, maintaining a database fully on enclave memory, where it cannot be accessed by anyone other than the CPU is another way to keep the data secure by leveraging isolated and trusted execution environments. Different techniques will be furthermore discussed below.

As we can see, isolated and trusted execution environments are an important feature when it comes to protecting the data from the OS and Key-Values Store systems do benefit from them.

2.3.2 Circumvention of SGX Limitations

There are a few limitations and challenges of the SGX platform that we address when programming for such technology.

It starts with a big challenge of choosing and defining what parts of the program can benefit of the SGX security. As it is known, it works with two major application components, the trusted and untrusted modules of our program. The limitations have to be thoroughly analysed so we can make that definition.

The main limitations are:

- Performance
- Memory
- I/O
- syscalls

In the KVS world, as we extensively covered, performance is the major concern and there is no real way around this limitation. Using secure enclaves will definitely decrease the supposed blazing fast performance. Although, with intelligent partition between the untrusted code, which will be fast, and the trusted instructions, which will be slower we can limit the performance overhead. By separating and well defining both modules of the application, we can decrease the code that needs to run securely and find a fine compromise between security and performance. Homomorphic encryption is also a mechanism to speed up performance by performing queries directly over encrypted memory. Although, fully homomorphic encryption is not a possibility yet, so, by compromising some set of operations that are not possible, partial homomorphic encryption will speed up the performance.

With the SGX base support the access to EPC from the owner enclave is efficiently processed by hardware en/decryption logics at cache-line granularity. This means that when the cache-line is brought from EPC to the processor, it is decrypted. The hardware logic calculates the keyed hash value of the cache-line, and verifies it against the stored hash value of the address. Internally the integrity hash values are organised in data-structures similar to Merkle Trees, to allow sub-trees to be evicted from the on-chip

storage securely [44]. Due to the space and time overhead of storing and processing security metadata at fine-grained cache-line granularity for EPC, the EPC capacity is unlikely to increase significantly. For example, a huge Merkle tree for tens gigabytes of main memory at cache-line granularity will necessarily increase the integrity verification latency intolerably, as an inefficient solution that will sacrifice throughput and latency [91].

Memory sizing is also a limitation when using enclaves in SGX technology. The amount of private secure data that can be maintained by the enclave is limited to the size of the enclave cache, which is around 128 MB, being that only about 94 MB are available to the application, with the rest reserved to metadata. Now, with SGX v2 and for some operation systems, mainly Linux because of paging swap support, it can be increased up to all the memory available in the system [87] by swapping pages from the EPC to main untrusted memory, with guaranteed of confidentiality, integrity and data freshness. When evicting pages from the EPC, it is assigned a unique version number which is recorded in a new type of EPC page and the contents of the page, metadata, and EPCM information are encrypted and written out to system memory. When reloading a page back into EPC the page is decrypted and has its version and integrity checked to make sure it was not tampered with.

Although, page eviction to main untrusted memory introduces a big overhead because of encryption and decryption and integrity checks (2x - 2000x) [14]. Clever partitioning of the application into the untrusted and trusted modules will help to overcome this limitation as described in the next sections. Another issue is EPC fault handling, because EPC limit requires exiting the enclave improving more the cost of paging.

I/O and syscalls are limited by default on the enclave for security purposes, so it can't affect or be affected by the OS. There is a way to perform and access I/O and syscalls through the aforementioned Ecalls and Ocalls (section 2.2.5 of this thesis), but they have to be accounted for when implementing the application. To address the problem, recent proposed solutions try to reduce the frequencies of enclave exits for system calls by running threads in untrusted execution for the interaction with the operating system, communicating with the enclave thread by sharing memory. Also, Ecalls and Ocalls require exiting and entering the enclave and that carries a big performance overhead, as encryption and decryption cycles must be done to maintain security guarantees, as well as some integrity checks.

2.3.3 SGX-Enabled Secure Databases

Database management service developers are now implementing secure databases ready to take advantage of Intel SGX hardware. It differs from normal databases because it runs on top of protected and encrypted memory so it can work with minimal TCB.

Next subsections present and discuss the overview of the current technology that leverages SGX to provide a secure database.

2.3.3.1 EnclaveDB

EnclaveDB [74] is a privacy enhanced and secure database that works alongside with Intel [SGX](#) and provides a Structured Query Language (SQL). It uses its technology to maintain **all** sensitive information inside [SGX](#) enclaves in order to keep them secure from a threat model of strong adversaries that can control the entire software stack on the database server. It resists attack from the administrator server, the database administrator and attacker who may compromised the operating system, the hypervisor or the database server.

Following Intel’s application guidelines, EnclaveDB has a two part architecture: trusted (running on the enclave) and untrusted modules. The enclave hosts a query processing engine, natively compiled stored procedures and a trusted kernel which provides API’s for sealing and remote attestation. The untrusted host process runs all other components of the database server. Figure 2.7 shows the architecture of the enclaveDB server-side.

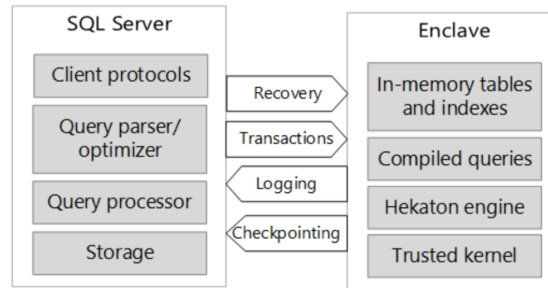


Figure 2.7: Server-side components of EnclaveDB

Leveraging [TEE](#), EnclaveDB then provides a database with a [SQL](#) interface and guarantees confidentiality and integrity with low overhead. With its design it also reduces the [TCB](#) to a smaller set than any other "normal" database.

2.3.3.2 Pesos DB

Pesos [61] is a secure implementation of object storage services like Amazon S3 [7], Azure Blob Storage [21], Google Cloud Storage [39] among others. In these current large-scale services, due to their complexity, the risk of confidentiality and integrity violations increase significantly. This storage systems are characterised by multiple layers of software and hardware stacked together which means the access policies for ensuring confidentiality and integrity are scattered across different code paths and configurations, thus exposing the data to more security vulnerabilities. Furthermore, untrusted third-party cloud platforms expose an additional risk of unauthorised data access by a malicious administrator.

Pesos allows clients to specify per-object security policies concisely and separately from the remaining storage stack. It also provides cryptographic attestation for the stored objects and their associated policies to verify the policy enforcement.

It enforces this policies by leveraging the Intel [SGX](#) for trusted execution environments and Kinetic Object Storage [60] for trusted storage (secure persistent storage - not the focus of this thesis). It structures a policy-compiler, its binary-format interpreter, per-object policy metadata, and the enforcement logic into a single layer of the storage stack. With this unification, it drastically reduces the [TCB](#) when compared to the order cloud services. Then it uses the trusted execution environment provided by [SGX](#) to connect directly Kinetic disk through an encrypted Ethernet connection allowing for object transfer and policy enforcement securely without any intermediate layers in the storage stack.

2.3.3.3 Speicher

Speicher [22] is a secure [LSM](#)-based (Log-Structured Merge Tree) Key-Value store that uses Intel [SGX](#) and it ensures not only strong confidentiality and integrity properties, but also data freshness to protect against rollback/forking attacks. It leverages [SGX](#) technology to achieve those security characteristics focusing on providing a **persistent** service, tolerant to system faults and securely recovering from crashes. It also tackles in interesting ways, two of the major limitations of [SGX](#): Memory Limits and Performance.

Implementing a Key-Value Store has a major requirement - High performance and low latency queries for big data structures. As already discussed, [SGX](#) has some memory limits. The enclave memory is located in the Enclave Page Cache ([EPC](#)) which is limited to 128 [MB](#) with about 94 [MB](#) available for application use (the rest being reserved for metadata). To allow creation of enclaves with bigger size than [EPC](#), the [OS](#) can use secure paging mechanism where it evicts pages to untrusted memory. Although with page encryption, decryption and integrity checks, this solution introduces high overheads (2× - 2000×) [14].

To address this performance and memory problems, the developers of Speicher implemented the following custom features (from Speicher public paper):

- **"I/O library for shielded execution:** Direct [I/O](#) library for shielded execution. The [I/O](#) library performs the [I/O](#) operations without exiting the secure enclave; thus it avoids expensive system calls on the data path."
- **"Asynchronous trusted monotonic counter:** Trusted counters to ensure data freshness. The counters leverage the lag in the sync operations in modern [KVS](#) to asynchronously update the counters. Thus, they overcome the limitations of the native [SGX](#) counters."
- **"Secure [LSM](#) data structure:** Secure [LSM](#) data structure that resides outside of the enclave memory while ensuring the integrity, confidentiality and freshness of the data. Thus, the [LSM](#) data structure overcomes the memory and [I/O](#) limitations of Intel [SGX](#)."

The technology leverages [SGX](#) with a clever partition between trusted and untrusted modules of the application. By maintaining the encrypted data on untrusted memory

hardware it addresses the memory and persistent limitations, and by keeping some information in secure enclave memory and with a good I/O library it overcomes (to an extent) the performance issues.

2.3.3.4 ShieldStore

ShieldStore [59] is a "(...) shielded in-memory Key-Value Storage with SGX". It aims to provide a very fast and low latency queries over very large data trying to overcome the SGX memory limitation. It accomplishes it by maintaining the majority of the data structures in the non-enclave memory region, addressing as well the performance issue by not relaying on the page-oriented enclave memory extension provided by SGX.

ShieldStore runs server-side in the enclave to protect encryption keys and for remote attestation and it is used to perform all the KVS logic. It uses a hashed index structured but places it in the unprotected memory region instead of the enclave EPC. As the main data structure is not protected by the SGX hardware, each data entry must be encrypted by ShieldStore in the enclave, and written to the main hash table.

The main flow and architecture is as described on figure 2.8.

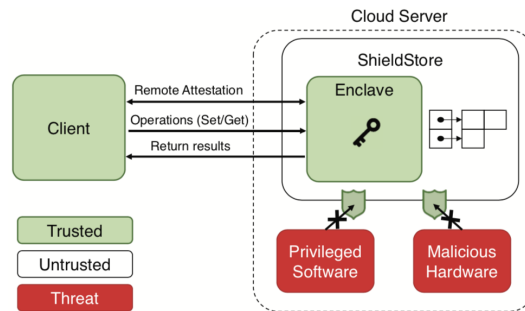


Figure 2.8: Overview of ShieldStore

First the client remote attests the server-side (1) verifying SGX support of the processor, the code, and other critical memory state of an enclave. In a second step, the client and the server exchange sessions keys (2) in order to establish a secure connection, using Intel SGX libraries to do so. Using this newly generated session key, the client sends a request for an operation (3). The server deciphers and verifies the request and accesses the Key-Value Store (4). Clients do not access the server-side ciphertexts neither need to know the encryption key used by the server to encrypt the values. The server will then decrypt the data from the storage, encrypted it again with the session key and reply to the client (5). All accesses to the KVS have integrity checks.

2.3.4 Discussion

Concluding, EnclaveDB (section 2.3.3.1) and Pesos (section 2.3.3.2) presents secure databases and objects storage systems respectively, using SGX, but EnclaveDB assumes that SGX

supports large enclaves whose size is an order of several hundred GBs and Pesos restricts the size of data structure to the size of EPC. On the other hand, Speicher (section 2.3.3.3) and ShieldStore (section 2.3.3.4) proposes a store that alleviates the memory limitation of Intel SGX by storing encrypted data on untrusted memory regions. Speicher and ShieldStore have similar architectures, but the former is primarily design for persistence storage and the latter is focused on a fast in-memory key-value store.

We can now conclude that the clever partitioning of the application into trusted and untrusted parts is really important when programming with Intel SGX. It directly affects syscalls and I/O, performance and memory of the service.

The long goal of this thesis is to implement a system with characteristics from the databases present above. In terms of performance we plan to implement partial homomorphic encryption, so it allows to perform operations directly over in-memory encrypted data. This will be a challenge, as fully homomorphic encryption is not yet practical [38], so adaptations must be made, but performance increase are expected over the databases presented above, by not needing to decrypt the data in secure execution. Persistence will also be a requirement just like some of the databases presented.

For trusted execution with SGX, extensive research is needed to partition the application in the two necessary modes to circumvent persistence, performance and memory SGX limitations. It will also be researched and tested the ability to provide built-in replication and availability with SGX.

2.4 SGX Virtualisation Frameworks

Virtualisation is the mechanism that relies on software to simulate hardware components and with the objective of creating a simulated ("virtual") computer system. It means that multiple virtual machines can run inside the same physical machine, and although sharing physical resources between them - cpu, memory, disks, run completely separated from one another.

With the broader adoption of virtualisation and SGX over the years, developers are working on ways to leverage virtualisation with the secure enclave technologies.

Also, container-based virtualisation has become very popular in the last few years. Container can use the capabilities of a secure and trusted environment to achieve a fast, highly portable, faster to deploy, and small SGX environment that can run mainly unmodified applications without a big performance overhead.

2.4.1 KVM-SGX

"KVM (for Kernel-based Virtual Machine) is a full virtualisation solution for Linux on x86 hardware containing virtualisation extensions (Intel VT or AMD-V). It consists of a loadable kernel module, `kvm.ko`, that provides the core virtualisation infrastructure and a processor specific module, `kvm-intel.ko` or `kvm-amd.ko`" [56].

The [KVM](#) framework allows to run multiple virtual machines with unmodified Linux or Windows images. Each [VM](#) will be assign their own virtual private resources like the network cards, disks graphic adapters, etc...

The KVM-SGX project is a module KVM/Linux module to support SGX virtualisation on [KVM](#). It exposes a private virtual SGX that can be used by the guest [OS](#). To fully virtualise [SGX](#), [EPC](#) must be assigned to the guest [OS](#). Although, since the [EPC](#) is a system resource, the management of this physical [EPC](#) - that backs the virtual [EPC](#), its owned by the kernel. The kernel exposes a device `/dev/sgx/virt_epc` that allows the guest [VM](#) to make a system call that creates and assigns its own virtual [EPC](#) section.

To recap, the guest [VM](#) does not have to deal with [PRM](#) or even [EPCM](#) because those resources are sitting below it on the hardware. The guest has a virtual Enclave Page Cache managed by the hardware that it can use at will to store secure code and data.

The Linux [SGX KVM](#) module is still in its development phase but can already be used to virtualise [SGX](#) platforms onto virtual machines [25, 50, 89].

2.4.2 Graphene-SGX

Graphene [42, 97, 98] is a guest OS design to run a single unmodified application with minimal host requirements. It bridges the gap of portability by abstracting the [OS](#) making it easier to port applications to different [OSs](#). Like a [VM](#), Graphene provides an isolated environment for application to run in, and, by not running a complete operating system it has a much lighter footprint, and generally low memory requirements and performance overheads. Also, developers are working to supply graphene as a containerised application [43].

As an extension, Graphene is working to provide integration with Intel's [SGX](#) processors where applications could run unmodified with the additional security guarantees of the secure enclave technology [97].

Since it is a library OS, Graphene takes advantage of its lightweight form factor to serve as a compatibility layer between the hardware and the application where normally, they do not work out of the box, with small performance overheads high portability and minimal or even no code modifications.

2.4.3 SCONE

The SCONE [14] (Secure CONTainer Environment) platform facilitates always encrypted execution using Intel's [SGX](#) secure enclaves to run encrypted code not even accessible to system administrators or root users. It supports this kind of execution running inside Docker containers and even managed Kubernetes clusters.

Although containers can have a big [TCB](#) and a weaker isolation, the design of scones leads to a small [TCB](#) by exposing a small C standard library that allows for system call to execute outside the enclave, but maintaining security guarantees by seamlessly and transparently encrypting/decrypting application data. Files stored outside of the enclave

are therefore encrypted, and network communication is protected by [TLS](#). SCONE also implements a *user-level* threading that maximizes the time that the thread spends inside the enclave. Combined with the asynchronous system where OS threads outside the enclave execute system calls, it minimises a big enclave overhead - **enclave entry/exit**.

2.4.4 Asylo

Asylo [16] is an open and flexible framework for developing enclave secured applications. Developed by Google, Asylo provides [17]:

- The ability to execute trusted workloads in an untrusted environment, inheriting the confidentiality and integrity guarantees from the security backend, i.e., the underlying enclave technology.
- Ready-to-use containers, an open source API, libraries, and tools so you can develop and run applications that use one or more enclaves.
- A choice of security backends.
- Portability of your application's source code across security backends

Asylo does not lock their technology with Intel's [SGX](#), but instead leaves it open to multiple secure enclave frameworks. In Asylo, the majority of user developed logic lives inside the enclave. However, due to security and portability reasons, this framework does not support direct interactions between the enclave and the [OS](#). All of enclave-to-OS interactions must be mediated through code that runs on the outside of the enclave but Asylo provides most of the code for creating, exiting and interacting with the enclave and the [OS](#).

2.4.5 Discussion

Virtualisation was a game changing technology when it was first created and, to this day, it's still evolving and it has become the core of cloud computing. With each iteration, more and more components can be virtualised and simulated to be used freely by their guest [OSs](#) and virtualised [SGX's](#) processors are starting to appear. The [KVM SGX](#) module and Graphene are a more low level types of virtualisation that can use [SGX](#) to provide a even more isolated and secure environment for application to run. However, with its fast shipping and deployment and highly portable applications, containerisation follows the same path, as it becomes more and more popular to developers and SCONE and the Asylo frameworks work beautifully to combine the power of a secure enclave with the characteristics of a container-based application.

In this thesis, we will use SCONE technology to run an unmodified Redis system secured with Intel's latest Software Guard Extensions ([SGX](#)) to expose a seamless secure application to the client with fast deployment speeds and high portability and scalability.

2.5 Related Work Balance and Critical Analysis

In the current days, computer scientists are always looking for a secure, fast and cheap environment to develop applications. As we know, it is not feasible to have all three of this elements working flawlessly without any compromises. Although, by combining in-memory key value stores, trusted remote execution environments and cloud providers, developers can now have a practical example of what would be to develop for a privacy enhanced secure system with reduce costs by using cloud providers and with better reassurances that the hardware and software that it's out of the control of the user will have a minimal impact on sensitive data and code of an application. By adding the performance benefits of an in-memory key value store and all of its technology, like built-in security, built-in replication and persistence we can in the best of our abilities today, combine the best of the three worlds without compromising too much on any of them.

In this thesis we will compare different kinds of approaches to implement a fast system with the assurance of a secure data flow that can easily be deployed into the cloud without fear of any components out of our control.

SYSTEM MODEL AND DESIGN OPTIONS

In this chapter we provide an overview of the system model and implemented architecture along with all of its components. It is explained how all component interact with each other and each component purpose and how they work.

It is also explained how the different security features are implemented into the system in order to provide a secure overall system and achieve the contributions planned for this thesis.

Section 3.1 provides a small recap and refinement of the objectives and contributions of this thesis and the used TEE technology.

In section 3.2 we describe the basic security assumptions need for our system to work securely. The threat model is a very important part of any security related project as it provides a clear overview of what attacks the system is able to protect against, presents the trustability chain and TCB as well as what components and security properties are out of scoped and not addressed in the project. Sections 3.3 and 3.4 presents an high level view of the system model and implemented architecture of the system. Then, in section 3.5 and 3.6 we expose the application specific operations supported both key value store specific operations and custom operations implemented by the proxy and also an interaction flow to represent a user interaction with the system.

Finally, as always, section 3.7 we summarise all the findings and provide a clear transition into chapter 4, which will present all implementation specific details performed to achieve the described system.

3.1 Refinement of Objectives and Contributions

The goal of this dissertation as explained on chapter 1, is the design, development and validation with experimental evaluation of a secure in-memory storage (based on a "key-value" model), supported by a hardware-enabled trust computing base.

Regarding the security assumptions, the solution will provide: (i) hardware-isolated in-memory processing engine, designed as a hardware-isolated container facility, enclaved within the Intel-SGX protection guarantees; (ii) hardware-isolated communication endpoints for client access, providing TLS tunnelling with strong TLS 1.3 endpoint encryption parameterisations and support for mutual client/server authentication, and (iii) privacy-enhanced operations to be directly processed on encrypted data sets in memory. The former facility is particularly interesting to combine the possibility to manage protected memory for small data sets and also searchable encrypted data sets that are far larger than the protected memory limits imposed by the SGX memory mapping facility. Furthermore, the solution will target main data structures commonly use fine-grained data items that can include pointers, complex composite types and keys, which do not match well with the coarse-grained paging of the SGX memory extension technique.

In the next subsections we will align the implementation ideas starting by refining the circumvention of limitations in SGX and the threat model assumptions to address our solution.

3.1.1 SGX Limitations Refinement

Section 2.3.2 describes in detail the limitations of SGX. In our approach we will intend to design a solution that can be leveraged from conventional reference KVS technology, giving the possibility to manage small datasets but also larger datasets (directly mapped in non-protected memory). To circumvent the problem our solution must combine the possibility to use the internal capabilities native to SGX with the possibility of supporting operations managing datasets encrypted in memory, with such operations executed directly over encrypted data. This facility will be provided by the use of partial homomorphic encryption constructions, with data initially encrypted and submitted to the key-value-store solution with cryptographic keys only managed in the client side. Our solution will be designed in order to be possible the support for fine-grained key-value encryption, driven from the application requirements. Our target is the support of a variety of operations provided in a typical KVS API, taking REDIS as the reference solution and in order to support a considerable number of queries currently used by many REDIS-supported applications.

3.2 Threat Model and Security Properties

The threat model and security properties definition describe the conditions of how we define a secure system. However, to ensure a basic secure system, we must achieve a few

key goals and objectives:

Data Privacy - Data must remain private to its owner.

Data Integrity - Data must not be compromised, modified or corrupted.

Authenticity - Data and system interactions should be authentic and not spoofed by unauthorised users.

Subsections 3.2.1, 3.2.2 and 3.2.3 explain under which conditions and assumptions the system will achieve these three main parameters and implement a secure system.

3.2.1 Adversarial Model Definition

As the baseline, our threat model will lie on the protection overview stated in SGX's paper [64]: *SGX prevents all other software from accessing the code and data located inside an enclave including system software and access from other enclaves. Attempts to modify an enclave's contents are detected and either prevented or execution is aborted*, which falls in the following adversary model:

Isolation by trusted containerisation from malicious code: The system performs and protects its data from an attacker capable of compromising the system through another application installed on the same system or malicious code existent or injected in the OS or OS hypervisor layers;

Privacy protection against insider "Honest but Curious" System Administrators: The system must be able to protect from an attacker with root access to the machine, with permissions to access and monitor memory-mapped data. This is relevant because we will target our solution as a candidate solution to protect data privacy in a cloud-based key value store solution as a service, preventing data-leakage vulnerabilities exploitable by insider incorrect users or system-administrators.

Network Attacks: All communication to and from the system (supporting client-service operations) should be secure, using proper strong cryptographic parameterisations for TLS 1.3, mutual authenticated handshakes and TLS endpoint executions isolated in SGX-enabled TLS tunnels in communication containers, avoiding attacks against the authentication of the service endpoints, as well as, attacks against the integrity and confidentiality of data flows supporting [REST/HTTPS](#) operations.

File system and memory access attacks: All sensitive data residing outside protected memory should be encrypted and operated in the encrypted form. An attacker can access the physical disks and hardware without the sensitive data being exposed.

3.2.2 System Assumptions

With the above security baseline considered for the threat model assumptions, the solution must be resilient to malicious privileged attacks and certain physical attacks. With a controlled and reduced hardware-shielded trust computing base, we want to design a solution that does not rely on the security of operating system managed by cloud providers. Furthermore, our solution must be also resilient to direct conventional physical attacks,

such as cold boot attacks, which attempt to retain the DRAM data by freezing the memory chip or even bus probing to sense and to read exposed memory channel between the processor and memory chips. The only weaknesses not covered in our concerns with be the SGX lack of protection for side-channel attacks.

The system planned has certain assumptions and aspects that are considered to be out of scope for this dissertation:

- **Trusted Client** - The client side is assumed to be completely trusted and correct.
- **DoS and DDoS** attacks are out of scope.
- **Side Channel Attacks** - It is out of scope any side channel attacks or any related attack not present in SGX's threat model.
- **Physical and Hardware attacks** exploring the SGX processing model and its isolation guarantees are out of scope, namely those presented above initially addressed in chapter 2, section 2.2.

3.2.3 Countermeasures for Privacy-Preservation

Under the described threat model and system assumptions, the system has some measures to achieved the desired security and trustability level.

All instances of the datastore are running in a containerised solution which means that each container is not only isolated from the host but they are also isolated from each other. Data is kept in memory at all times unless persistent disk storage is turned on. Data in main memory is secure when running both in protected and unprotected mode.

Privacy, integrity and authenticity when running in protected mode are always ensured by the SGX's technology under their threat model and assumptions. But when the system is running in unprotected memory, outside the trusted execution environment, the system will always keep data always encrypted with strong and standard state of the art cryptography algorithms, therefor preserving privacy of data.

For data integrity preservation, all values are appended with standard checksums calculations and integrity check algorithms. Authenticity is preserved by performing strong and standard cryptographic signing algorithms.

As for the application security, the system provides a secured API that only allows authenticated requests and it contains a role based segmentation authorisation with secure and strong passwords and cryptographic keys. This system also applies the principle of least privilege [63] to all actions, where users never have more privileges that they require.

Privacy is also preserved on all communications with the use of the strongest transport layer security algorithms currently available, and established trust between all components with a trusted certificate chain.

3.3 System Model

The main goal of this project is modelled in figure 3.1. As shown, the model can be divided in four main components - the **client**, the **proxy server**, the **key-value storage server** and the **authentication server**.




Figure 3.1: System Model Overview

The system complies with the adversary and threat model explained in section 3.2.1. The user is not aware of any implementation details and has seamless interaction with the system despite the architecture and implementation provided by the backend, meaning that all solution expose the same [API](#) and support an interface as equal as possible to the unprotected version of the key-value store.

All of the four main components will be explained in the subsections below.

3.3.1 Key-Value Storage Server

The storage server is a key-value store meant to hold data required by the user. This data is kept in memory so it allows fast read, writes, updates and deletes.

In this project, there will be two different kinds of storage servers. The **secured** and **unsecured** storage server. This nomenclature do not reflect the privacy, integrity and authenticity security problems as they are preserved on both components but it reflects the environment which they will run on.

The **secure** storage runs on a trusted execution environment (trusted hardware) protected by a secure physical processor present on the host machine. This processor provides the security features that allow for memory to be operated in plain text without breaking privacy or data security as explained in section 2.2.

As for the **unsecure** storage, it describes a storage service that runs on unprotected memory regions and on untrusted hardware. This means that data must be actively secure by encryption, integrity checks and authentication features in order to provide and preserve privacy.

We must mention that the above system model addresses the replication facilities as leveraged from the Redis solution and tries to offer the same availability conditions as the original Redis service. Both secure and unsecure configurations will provide availability by running replicated Redis instances in a cluster. This availability provided by the replication of Redis instances that can run as cloud service instances (that can be distributed among different cloud providers and geo-replicated datacenters) only offers consistency guarantees under a fail-stop model (not extended to byzantine fault tolerance or byzantine intrusion tolerance, that are conditions out of scope of our planned dissertation).

Every Redis instance running on a secure configuration will also provide an API for remote attestation, going forward known as the attestation service, so their hardware stack and software can be attested by remote parties, to provide trust to the users.

3.3.2 Proxy Server

The proxy server serves as a gateway to the storage instances. It is a single point of entry to the system which has some advantages. Not only it provides an abstraction to a backend which can be replicated across multiple instances across multiple geo-locations, it also can add additional features to the system.

With a centralised gateway we have a centralised access management with authentication and authorisation features. Not only that, but the connection to multiple instances of secure data storages with different authentication mechanisms and secrets can also be centralised in one place, which allows the user to have just a single username/password pair to connect to all the storage instances.

The proxy server manages security properties when connecting to instances on a unsecure storage configuration. It is the proxy that handles security features like encryption and decryption, integrity and authentication checks, and provides a completely encrypted and private storage service.

It also exports a [REST API](#) which can be used by multiple clients with different implementations and also offers custom features that the standard data storage system does not. Custom operations will be explained below, on section 3.5.

The proxy server serves as a single point of access for the multiple storage instances remote attestation features. It attests every instance and feedback the user on the hardware stack and software state of each instance. Furthermore, the proxy server itself runs on a trusted execution environment leveraging secure trusted hardware and can also be attested by the clients.

3.3.3 Authentication Server

An external authentication server is required to centralised user management. It is responsible for user authentication and verification. A user authenticates against the authentication server receiving an access token, that it provides to the proxy on every request. The proxy can verify the token and check if it is valid and what permissions this user has, and can authorise or reject the access to a particular endpoint or system functionality.

An external authentication server relieves the proxy of user authentication and management, and outsources it to a standardised open source system that implements security standards of identity and access management. This is also import so that the proxy server can, if needed, be replicated.

3.3.4 Client

The client, or user, is the one that will consume the exposed [APIs](#) by the proxy server. In the case of this project, the client will be a representative benchmark tester that will use the exposed endpoints to record benchmark times and various relevant evaluation criteria so all storage configurations and replication mechanism can be compared.

Although just a simulated client/tester, the [APIs](#) are consumed the same way a real user would consumed them, so, benchmarks are a representation of real system usage.

3.4 System Architecture

Figure [3.2](#) describes the hardware and software representative stack of the infrastructure ¹. Figure [3.2a](#) shows how the storage service are deployed onto the cloud provider's machines. The machines hardware provide a very specific and physical processor that implements a trusted execution environment. Running on the operating system, a containerisation solution runs the [TEE](#) virtualisation framework in order for the container running both the key value stores and the attestation services being able to access the [TEE](#). On the right hand side of this figure, we can see a storage service that runs with no [TEE](#) virtualisation system, and that is the deployment of an unprotected key value store configuration.

On figure [3.2b](#) the stack provided describes a system also running in the confinements of a trusted execution environment and therefor, also providing an attestation service.

3.5 Supported Operations

The exposed [API](#) can perform some operations over the storage system. Some operations are out of the box key-value stores operations, and some were implemented in the proxy so to give a customised and more complex operations.

¹The described stack is a representative stack of the machine layers and it does not represent a true overhead ladder




Figure 3.2: System Architecture Stack

The next subsections will iterate and explain the supported operations, and we need to indicate that all operations listed can work on both main data storage configurations (protected and unprotected), where operations are performed whether data is maintained on clear text, securely inside a trusted execution environment or maintained encrypted in unprotected memory regions.

3.5.1 Role-Based Authorisation

The system by default rejects all unauthenticated requests, which means, users do need to be registered and have credentials to access the system.

As explained before, user management is performed by the external authentication server, and each user has a role assigned. Different roles can do different actions on the system and role based authorisation can be completely customised.

In this thesis we operate with the principle of least privileges, and came up with two different roles: a **BasicUser** which can only perform read operations and cannot alter the system state (basically a readonly user), and an **Administrator** which can perform all actions on the system. These roles are just a representation of the role based authorisation that the system performs before accepting a request into the server.

3.5.2 Key-Value Storage Operations

The system is ready to perform a set of operations from the key value storage of choosing, without modification. However, this is not a full fledged solution, and accommodations had to be made for the tow main configurations of the storage system: protected and unprotected.

The supported operations are as follow:

- **Set** - Stores a value into the database associated with a key.
- **Get** - Retrieves a value associated to the provided key.

- **Set on List** - Creates a list of values associated to a key.
- **Get List** - Retrieves all values associated with the given key.
- **Set on List with Score** - Creates a list of values associated to a key. Each value has associated a score (integer)
- **Get on List between Score** - Retrieves all values associated with the given key with scores between provided scores.

3.5.3 Proxy Enabled Operations

The proxy server implementation allows for the implementation of another set of operations not supported out of the box from the key value store:

Sum, given a key and a number, the server can fetch the value associated with given key and add it to the number provided.

Subtraction, given a key and a number, the server can fetch the value associated with given key and subtract it to the number provided.

Multiplication, given a key and a number, the server can fetch the value associated with given key and multiply it to the number provided.

Search on List, where the server takes a search term, match it against each value of a provided key of a list and return only the matching values.

3.5.4 Attestation

Attestation is the method provided by the processor implementing a trusted execution environment to establish trustability assumptions to a remote challenger. It proves to a third party that the system is running the correct application with the correct system configurations in the expected physical hardware. This thesis implements two different kinds of attestation: An hardware and [TEE](#) attestation and an on-demand software and system stack attestation.

Figure [3.3](#) describes the [TEE](#) hardware attestation infrastructure. It happens at application start up, and is the attestation that relies on the secure processor and it requires an external Configuration and Attestation Server ([CAS](#)) that holds some kind of secret necessary for the application startup. However, the [CAS](#) will only allow and provide the secret if the application enclave can prove itself to be correct, i.e, is running the correct source code in the correct system. The application contacts a Local Attestation Server ([LAS](#)) to obtain a proof of their correctness and then try to fetch their secrets in the configuration server. If the application is deemed correct, it can access those secrets and can start.

On the other hand, users can also request the remote attestation of the complete system software stack and hardware information at any point on the user's lifecycle. The proxy server is responsible to contact each storage server's attestation service and request




Figure 3.3: Attestation Model

attestation quotes. Each quote is then gathered in the proxy, which attests it self and returned to the client.

The correctness of the system is then determined by the client which analyses the quotes provided by each system, and decides whether or not the system is correct and protected and can continue normal use.

3.6 Operation Flow

Figure 3.4 shows an example of a flow that can occur between all system components.

Following the flow, we can see that the first interaction is with the external authentication server by providing the user credentials (1) and (if login successful) receiving an access token in return (2).

Then, the user uses the access token to attest the system and make sure that hardware and software stacks are working as expected. It requests an attestation to the proxy (3) with the access token retrieved on login. The proxy checks the access token with the authentication server (4) and verifies if it is correct and has the necessary roles to access the attestation endpoint. If the verification fails, a proxy returns an authentication/authorisation fail back to the client (5).

If it succeeds, the attestation can proceed and the proxy server contact all storage server attestation services and gathers all quotes (6), and while it's doing that, also starts an self attestation process (7) and returns all quotes to the client.

Now, the client must analyse the quotes from the attestation and decide whether or not it trusts the system to continue (8). If it deems the system incorrect, it fails but if the client trusts the system, it can carry on with normal operations to the proxy and to the storage server (9).




Figure 3.4: Operation Flow

3.7 Summary

To achieve the objectives and contributions of this thesis, the system should rely on a physical hardware processor that can implement an isolated execution environment, to protect the system against insider threats, operating system and hypervisors vulnerabilities, and several other attacks already explained in the adversary and threat model of this project. That processor, as shown on system architecture section is present in the used machines, and to allow for a light weight, easily portable and deployable system, an TEE virtualisation technology was used in order to a containerised application can take advantage of the physical processor.

However this technology still has some limitations, and to address a memory limitations, the system also allows for a storage system to be placed in unprotected memory and still maintain security, privacy and integrity properties.

This way, we can maintain a system, running an unmodified key value storage application and deploy it safely on the cloud, or any other machine provider and be sure that data is as protected as it can be.

Not only that, but all the hardware stack and software can be attested so that, on a unlikely event of attack, data corruption or leaked vulnerability occurs, the user can always, to a high percentage of trust, define whether or not the underlying system is

correct and can be trusted.

PROTOTYPE IMPLEMENTATION

This chapter presents a detailed explanation of the implementation of the prototype - A Trusted and Privacy-Enhanced In-Memory Data Store, and all the implementation details that helped the system to achieve a secure state according to the adversary model.

Section 4.1 explains the system model presented on figure 3.1 from a developer view, and presents all used technologies, programming languages and implementation details used to achieve the desired system.

Section 4.2 presents some general additional security and implementation features also worth mentioning and in section 4.3 it is explained some tradeoffs decided in the implementation of the prototype, and why where they made.

To finalise, there is a general summary of the chapter in section 4.4 that gathers all important implementation features from all components.

The implemented prototype source code is available publicly on GitHub, secure data-store [85], the proxy [86] and the client/tester [84].

4.1 Architecture and Implementation Options

To achieve the goal of deploying the system in a cloud, we had to find a provider that has and provides host machines with the pretended TEE technology - Intel's Software Guard Extensions (SGX) v2.11. Although not globally available, some cloud providers are starting to make them available and for this thesis, the cloud provider used is OVH Cloud [73].

For this thesis, OVH provided an IaaS stack machine running Ubuntu Server version 18.04 with kernel 4.15.0-101-generic, which means that we have control over all host's stack but the hardware, from the operating system, networks, runtime and applications. The used machine specific configurations are listed on listing 4.1.

Listing 4.1: Machine Specifications

Dedicated Server Node
Processor: Intel 2x Xeon Silver 4214 - 24c/48t - 2.2GHz/3.2Ghz
Memory: 192 GB
Hard Drive: NVMe, SATA available
Public Network: Beginning at 1 Gbps
Private Network: Beginning at 2 Gbps
CloudLinux (Ubuntu 18.04 LTS Server 64 bits)

This particular Intel processor offers [SGX](#) with an [128MB](#) of enclave page cache (EPC) with about [94MB](#) being available for application use like explained in section [2.3.2](#) and all [SGX](#) linux drivers and [SDKs](#) were installed [[52](#), [62](#)].

All components of the application will be deployed using Docker v19.03.6 [[32](#)] and the Docker Compose tool v1.17.1 [[33](#)]. To integrate and run unmodified applications with [SGX](#), the SCONE v4.2.1 [[14](#)] technology was used, and will wrap all components that need to run within a secure and isolated environment.

4.1.1 Secure Redis

Redis [[76](#)] is the key-value storage server used by this thesis. Redis instances will run in two different modes, as explained in section [3.3.1](#). Unsecure Redis configuration will run on unprotected memory on a docker image based on the official Redis Docker repository [[34](#)]. For the secure configuration, SCONE framework already provides a curated image from their repository which contains a Redis server version 6.0-rc1 ready to run on an isolated environment, in this case Intel's SGX. The SCONE version used is the SCONE 4.2.1 to match across all the SCONE components.

Although all redis servers run behind a proxy all the necessary security features provided natively by the server are used. Only communications incoming from the proxy server are allowed and all are encrypted with strong [TLS 1.3](#) ¹ protocols with enclave termination. The non encrypted communication port is disabled, and mutual [TLS](#) authentication is turned on, which means that all clients are required to provide a certificate signed by the thesis CA in order to establish a connection.

Access Control is also enabled through an explicit [ACL](#) ². Following the principle of least privilege, users are defined via an username and a strong password and have permissions to access only the operations that they require to function.

When running in a replicated environment, master-slave or cluster, the same principles apply. Communication between replicas is also always through mutual [TLS](#) authentication, even in cluster mode where an event bus is necessary for replica synchronisation. Replicas are read-only and since they can connect to the master instance, they use a specific user with permissions to perform just the operations that the replica needs to synchronise, and cannot alter the state of the master instance.

¹TLS is a new feature released in Redis v6.0

²Redis ACL is a new feature released in Redis v6.0

4.1.2 Proxy Server

The proxy server is the component that abstracts the Redis configurations in the backend. Proxy is a spring boot starter, version 2.3.0.RELEASE, web server application written in Kotlin v1.4.10, a language that runs on the JVM with Java OpenJDK version 11.0.9.

This component serves as a single point of entry to the system, and clients connect to it via an exposed HTTP API via a SSL/TLS encrypted channel (HTTPS). This connection only authenticates the server, but users need to provide an authorisation header to access the server. The bearer token, on the format of a JWT (Json Web Token) must be provided by the external authentication server as the proxy will check with it to validate the request. Not all users can access all endpoints, and the proxy decides the access control based on the role presented in the token.

Using an external configuration file, the proxy is able to communicate with multiple configurations of Redis instances. Communication with the instance is performed via Jedis v3.3.0 [55], a simple and lightweight java Redis client. When connecting to a protected Redis instance, an instance secured by SGX processor and running inside an enclave, the proxy passes through the keys and values to the instance without any modification, meaning that values are secured inside the enclave even though they are handled in plaintext. However, all data residing on unprotected memory should be encrypted and by enabling a flag in the configuration file, the proxy will encrypt, sign and perform integrity checks on all keys and values. The homomorphic encryption is performed with the help of the Hlib v1.2r2 [90], an homomorphic encryption library implemented by the NOVA LINC developers.

Keys for all value formats (simple, lists, etc..) are always encrypted with the Homomorphic Deterministic (HomoDet) cipher, that guarantees the same encrypted string for the same clear text value. This allows for Redis to match a given key with one present in the storage without revealing the actual value of the key.

The value from the key-value pair is encrypted in a more complex way than the keys, and their format is detailed in 4.1, 4.2 and 4.3.

$$EncryptedValue = [value]_{K_s} \quad (4.1)$$

$$CompositeValue = [value]_{K_s} \mid [EncryptedValue]_{K_{signature}} \quad (4.2)$$

$$[value]_{K_s} \mid [EncryptedValue]_{K_{signature}} \mid [CompositeValue]_{K_{Hmac}} \quad (4.3)$$

The value is encrypted in one of two ways - referring to 4.1:

- **If the value is a string**, it is encrypted with a strong AES cipher working on a ECB (Electronic Code Book) mode with a PKCS5 Padding (AES/ECB/PKCS5Padding) from SunJCE provider, with a 256 bit key.
- **If the value is an integer/long/double**, the value is encrypted with the Homomorphic Addition (HomoAdd) cipher from the Hlib library with a Paillier Key.

Strings are encrypted with the strongest cipher because no operation will be performed over the encrypted value, but for arithmetic values, the `HomoAdd` cipher is used, to allow for addition, subtraction and multiplication operations over the encrypted value.

Regardless of the encryption cipher, the encrypted value is signed with a standard SHA512 with `RSA` signature algorithm from the `SunRsaSign` provider - 4.2. The signature is performed over the encrypted value to allow the homomorphic operations results being sign without having to decrypt the value. The signature is then appended to the encrypted value and the result is hashed with a `HMdSHA256` algorithm to provide a rapid integrity check. The result of all security operations, encryption, signature and hashing is appended into one string and set into the Redis database as a single value.

When setting a value in a list with a score, the keys and values are secured like explained above, but the scores can also pretend to sensitive data and are also encrypted. However, to provide the capabilities of fetching values between scores, a score is encrypted with the `HomoOpelnt` cipher from `Hlib`, an order preserving cipher. This cipher keeps the scores confidential but still preserves order and can be searchable on an inequality operation, for example $x < score < y$, that fetches all values with scores between an x and a y number.

For demonstration purposes, when adding a value to a set, on the respective endpoint, the value is encrypted with the `HomoSearch` cipher from `Hlib`. This cipher allows searching an encrypted value for a specific substring provided by the client.

The [API](#), endpoints and parameters, are completely documented and available on an OpenAPI v3.0.0 *yml* format.

4.1.3 Client-based Benchmarks

As explained before, the client is going to be emulated by a tester. Benchmarks were performed in two different ways: directly against the Redis instances, using the **redis-benchmark** [47] tool, and through the proxy [API](#).

This proxy [API](#) tester is implemented using the Gatling v3.2.1 [37] framework. Gatling is a load and performance testing tool that is configured as code. The configuration of the tool is written in Scala v2.12.3 and provides a very configurable [API](#). It has the ability to generate fields, that will be used to populate the database, and perform as much requests as needed or perform request during a certain amount of time. This framework also allows to perform the necessary login request to the external authentication server in order to provided the access token to the proxy [API](#).

For performance and load testing we can also make various simultaneous users perform actions at the same time and set up ramp up periods of higher load.

The benchmarks are doing requests to every endpoint available, and will compared against other proxy and database configurations and Gatling provides a detailed report for each one, exposing different metrics that will be presented further down this document in chapter 5.

Also, the same tests were configured on the Apache Jmeter v5.3[9] load testing platform to corroborate some results.

4.1.4 Authentication Server

The external authentication server is Keycloak v10.0.2 [57], open source identity and access management from Red Hat. This service implements the current standards for single sign-on, authentication and authorisation security.

Communication is performed over encrypted [HTTPS](#) channels, and the client logs in directly against this server to obtain the access token necessary to access the proxy.

User and roles management is done via the Keycloak user console and the platform allows for the configuration of different access token signing keys, their expiration date and supports token revocation and key rotation.

4.1.5 Attestation

Attestation is the mechanism responsible to provide the user with a trusted indication of the complete system state and is performed with two different methods, in two different scenarios. The hardware and [SGX](#) attestation and an on-demand software and system stack attestation

SGX attestation is performed at startup and its transparent and automatic using physical [SGX EPID](#)-based attestation between the application, a locally deployed attestation service (SCONE [LAS](#)) and the remote SCONE [CAS](#).

The SCONE Configuration and Attestation Service ([CAS](#)) is a part of the system's infrastructure and it is meant to hold application secrets and injected into authenticated and attested enclaves. This service is also running inside an enclave so, any secrets stored in the server are also isolated and protected from outsiders. Each secret has a strong access policy attached and guarantees that only the right enclave, running the correct code in the correct system can access them. On application start-up, the created enclave communicates with the [CAS](#) to fetch the necessary secrets for the application to start. This is the first step of the application process and it is described in figure 4.1.

Having established a connection with [CAS](#) (1), the application now contacts the SCONE Local Attestation Service in order to receive a signed quote from the trusted [LAS](#) (2, 3, 4) that then sends to the [CAS](#) (5). This quote contains information about the application enclave signed by the trusted [LAS](#), including and most important, the enclave hash, an hash determined by content of the pages of the enclave. [CAS](#) can now verify the quote provided (6) and return a success message (7) if the quote is deemed correct and signed by [LAS](#). Then, the application requests the secrets from [CAS](#) (8) and if the request passes all access policies (9), the secret is then returned (10) and injected into the enclave. These secrets can be injected as an environment variable or a file in a provided path, but whatever the format, they always stay inside the enclave protected memory, which means that no human or process, regardless of their roles in the system will be able to access

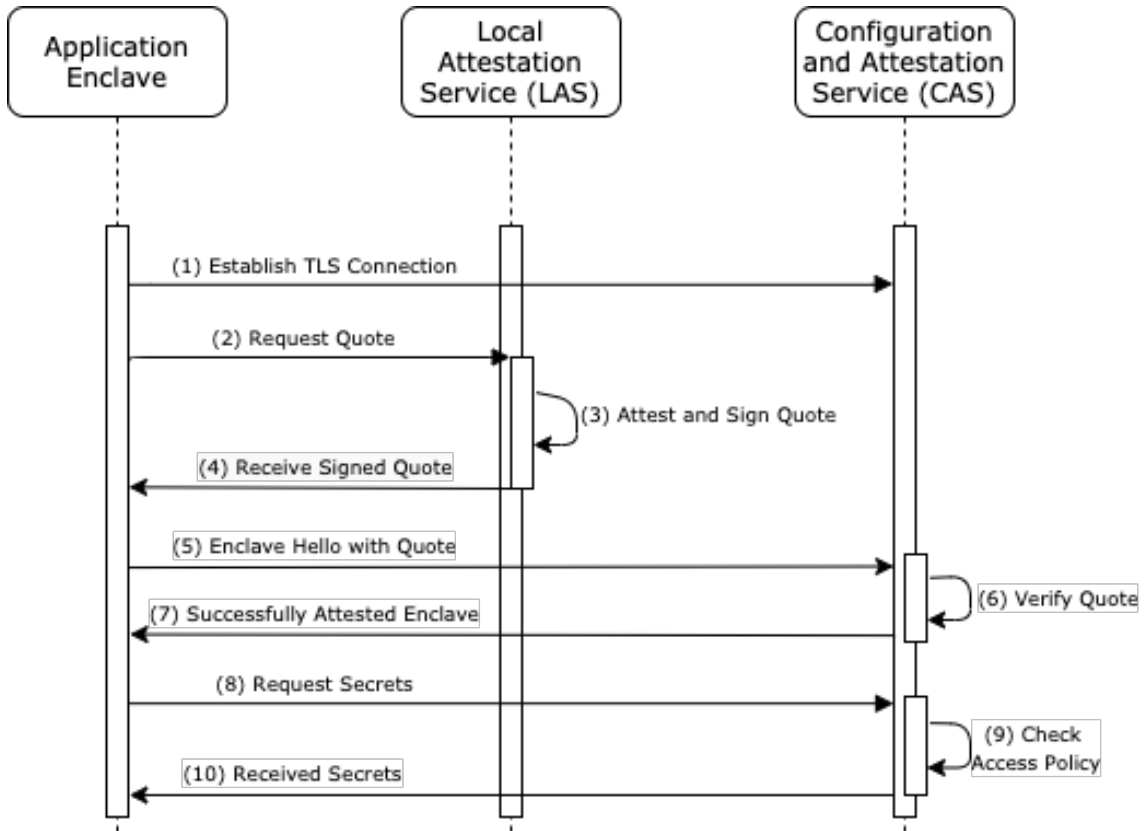


Figure 4.1: Attestation Flow

them. on enclave destruction, the secrets are also destroyed and never written to disk or persisted anywhere.

In case of the Redis storage server, the [TLS](#) private keys and certificates are stored in [CAS](#) and not in the container. This guarantees that, if it is able to provide the right certificates on a [HTTPS](#) connection, the enclave was attested and its correct. For the Redis custom attestation service, the secrets hidden in [CAS](#) are the quote signing keys and the same principle applies.

The application software and stack attestation relies on the first type attestation in order to access the keys necessary to operate, and it is performed at user request. This attestation is meant to attest binaries, config files, hardware and [OS](#) information to the client by providing signed hashes of those resources. The client then decides if they match and pass the expected integrity checks, and if sign with a key provided by the [CAS](#).

This attestation service was implemented in C++14 for the Redis instances, compiled using *Linux Musl* g++ GNU Compiler Collection (GCC) version 10.2.0 using the Linux Chilkat C/C++ Library v9.5.0.84 statically linked to the code in order to run totally inside the enclave, and Kotlin for the proxy server.

4.2 Additional Details

4.2.1 Protected Memory Check

SCONE framework says that it is running the applications inside secure enclaves and application memory is protected and there are two ways to verify it. By running an application with secrets in the SCONE configuration and attestation server ([CAS](#)), we can be sure that the program is running inside an enclave since CAS requires a specific [SGX](#) attestation mechanism.

On a lower architecture level we can actually inspect the memory being used by the application. To test this, we can write a simple C program that holds a secret on an array in memory and dumping the memory of the process via the `/proc` filesystem. The `/proc/<pid>/map` shows the different memory regions of the process and `/proc/<pid>/mem` holds the memory. By analysing this path, we can check that the secret is **not** present, as opposed to the same program running outside protected memory [\[36\]](#).

4.2.2 TLS, HTTPS and Certificate Chain

All communications in the system, both from outside or inside the system are performed over [HTTPS](#), [TLS](#) v1.2 or v1.3, with custom and different certificates for each component. However, we can establish a chain of trust by signing all certificates with a root custom certificate authority ([CA](#)). A custom [CA](#) was created, and signed all certificates and keys and marked as a trusted certificate authority for the entire system.

4.2.3 Logging and Auditing

All operations in the system, being logins, proxy standard operations or even attestation requests are logged. A log line contains a timestamp, remote [IPv4](#) address, request information such as path, method, and response status, and the request owner, the user. However, to protect preserve privacy, username is hashed. A system administrator can access the logs in order to audit the system.

4.3 Tradeoffs on the Implementation Options

Tradeoff is the loss a system property is exchange of another. When implementing additional security properties in a system, there will always be a performance impact. It is then up to testing and use case evaluation to decide whether or not the security increase compensates the performance decrease.

When it comes to replication, there is also a decision to make referring to a consistency-performance tradeoff. Being a key-value store, Redis is always focused on performance and that is why, it follows an eventual consistency model on replication, meaning that a Redis master node will propagate changes to their slave nodes but will not wait for slave acknowledge to respond to the client.

SGX technology also makes several tradeoff between performance and security, although it default to the latter. Protected memory size is a big issue to an **SGX** enabled application, because, as explained in section 2.3.2, the protected memory is limited, and when it is exceeded, protected memory is swapped to main memory. With additional encryption and decryption cycles as well as integrity checks, that as a penalty performance that can reach over 2000x.

To explain this problem, we can use the OpenSSL library as an example: Dynamically linking a library to an enclave will incur a security level penalty because, not only the library cannot be trusted as it is part of the operating system and can be compromised, but also **TLS** termination and decryption would be performed, outside the enclave. On the other hand, statically linking the OpenSSL libraries (option taken on this project) means that the enclave has the necessary libraries to perform **TLS** termination inside the enclave. However, it will increase the size of the application, reducing the amount of memory that can be present inside the enclave and having to be being swapped into main memory.

Another implementation detail on the performance-security tradeoff topic are the homomorphic ciphers that the proxy implements on an unprotected Redis instance. Homomorphic encrypted values do leak some minor information to an attacker, although it is a good trade-off between security and performance. For example, order preserving ciphers although they do not leak specific values, they do leak the decrypted value order.

The security-performance tradeoff is not a static line, and it moves depending on the use case. For some applications, data must be extremely secure and performance can take a hit to maintain a strong privacy preserving system. This thesis implemented various security properties and all are evaluated in chapter 5 to provide enough information for a user to decide whether or not it is a right solution to its own use case.

4.4 Summary

To recap, Redis will server as a storage server and runs on an isolated environment using Intel's **SGX** secure processor and also runs on normal unprotected memory. However, both solutions preserve privacy and integrity, either by **SGX**'s guarantees or by proxy enabled encryption and security properties. The proxy enables some Redis out-of-the-box operations, but also some additional operations. When Redis is running in unprotected mode, the proxy enables the performance of operations over encrypted memory, sparing encryption and decryption cycles.

Both Redis and the Proxy Server are attested at start-up, by requesting application secrets hold by a third party server that implements strong access policies to make sure each secret is injected to the correct enclave.

Users can also request attestation of the system on-demand, by requesting it to the Proxy server. The Proxy contacts all redis instance's attestation services, which are deployed alongside Redis, and returns binary, config files and system information that the

user can analyse and have the guarantee that it is communicating with a correct system.

In regards to communication, all channels between all components are secure with the standard string TLS encryption with a system wide CA signature.

These security properties come with some tradeoffs but they all address a necessary privacy property, and the next chapter will evaluate them and determine if there is a bearable overhead in order to implement a secure correct and trustable system.

VALIDATION AND EXPERIMENTAL EVALUATION

5.1 Testbench Environment

5.1.1 Cloud-Based SGX-enabled REDIS Single Instance

5.1.2 Cloud-Based SGX-Enabled REDIS Cluster

5.2 Benchmarks

5.3 Performance evaluation for Testbench 1

5.3.1 Latency

5.3.2 Throughput

5.4 Performance evaluation for Testbench 2

5.4.1 Latency

5.4.2 Throughput

5.5 Evaluation of the Attestation Protocol

Get startup times with and without CAS, get stack attestation times...

5.6 Complementary Measurements

Memory, CPU, etc instrumentation, workload

5.7 Summary and Findings

CONCLUSIONS

- 6.1 Main conclusions**
- 6.2 Open Remarks**
- 6.3 Future Work Directions**

BIBLIOGRAPHY

- [1] *80% Of Enterprise IT Will Move To The Cloud By 2025*. Accessed: 2020-08-12. URL: <https://www.forbes.com/sites/oracle/2019/02/07/prediction-80-of-enterprise-it-will-move-to-the-cloud-by-2025/>.
- [2] *A brief history of cloud computing*. Accessed: 2020-01-11. URL: <https://www.ibm.com/blogs/cloud-computing/2014/03/18/a-brief-history-of-cloud-computing-3/>.
- [3] *A history of cloud computing*. Accessed: 2020-01-11. URL: <https://www.computerweekly.com/feature/A-history-of-cloud-computing>.
- [4] *Aerospike*. Accessed: 2019-06-16. URL: <https://www.aerospike.com>.
- [5] *Aerospike*. Accessed: 2019-06-16. URL: <https://www.aerospike.com/docs/guide/security/index.html>.
- [6] *Amazon Dynamo DB*. Accessed: 2019-06-16. URL: aws.amazon.com/dynamodb.
- [7] *Amazon S3 - Cloud Storage*. Accessed: 2019-07-09. URL: <https://aws.amazon.com/s3/>.
- [8] *Amazon Web Services*. Accessed: 2019-06-16. URL: <https://aws.amazon.com>.
- [9] *Apache JMeter™*. URL: <https://jmeter.apache.org>.
- [10] *Apple's iCloud service suffers cyber-attack in China, putting passwords in peril*. Accessed: 2020-01-11. URL: <https://www.washingtonpost.com/news/the-switch/wp/2014/10/21/apples-icloud-service-suffers-cyber-attack-in-china-putting-passwords-in-peril/>.
- [11] *ARM TrustZone*. Accessed: 2020-01-08. URL: <https://developer.arm.com/ip-products/security-ip/trustzone>.
- [12] *ARM TrustZone Security Whitepaper*. Tech. rep. PRD29-GENC-009492C. ARM Limited, Dec. 2008. URL: http://infocenter.arm.com/help/topic/com.arm.doc.prd29-genc-009492c/PRD29-GENC-009492C_trustzone_security_whitepaper.pdf.
- [13] *ARM TrustZone Stack Image*. Accessed: 2020-01-08. URL: <https://malware.news/t/introduction-to-trusted-execution-environment-arms-trustzone/20823>.

- [14] S. Arnautov, B. Trach, F. Gregor, T. Knauth, A. Martin, C. Priebe, J. Lind, D. Muthukumaran, D. O’Keeffe, M. L. Stillwell, D. Goltzsche, D. Eyers, R. Kapitza, P. Pietzuch, and C. Fetzer. “SCONE: Secure Linux Containers with Intel SGX.” In: *12th USENIX Symposium on Operating Systems Design and Implementation (OSDI 16)*. Savannah, GA: USENIX Association, 2016, pp. 689–703. ISBN: 978-1-931971-33-1. URL: <https://www.usenix.org/conference/osdi16/technical-sessions/presentation/arnautov>.
- [15] *Around 75% of Open Redis Servers Are Infected With Malware*. Accessed: 2020-02-18. URL: <https://www.bleepingcomputer.com/news/security/around-75-percent-of-open-redis-servers-are-infected-with-malware/>.
- [16] *asylo*. Accessed: 2020-09-20. URL: <https://asylo.dev>.
- [17] *asylo*. Accessed: 2020-09-20. URL: <https://github.com/google/asylo>.
- [18] *Attestation and Trusted Computing - CSEP 590: Practical Aspects of Modern Cryptography*. Accessed: 2019-11-26. URL: <https://courses.cs.washington.edu/courses/csep590/06wi/finalprojects/bare.pdf>.
- [19] *Attestation Identity Key (AIK) Certificate Enrollment Specification*. Accessed: 2019-11-26. URL: <https://www.trustedcomputinggroup.org/wp-content/uploads/IWG-AIK-CMC-enrollment-FAQ.pdf>.
- [20] *AWS hit by major DDoS attack*. Accessed: 2020-01-11. URL: <https://www.techradar.com/news/aws-hit-by-major-ddos-attack>.
- [21] *Azure Blob Storage*. Accessed: 2019-07-09. URL: <https://azure.microsoft.com/en-us/services/storage/blobs/>.
- [22] M. Bailleu, J. Thalheim, P. Bhatotia, C. Fetzer, M. Honda, and K. Vaswani. “SPEICHER: Securing LSM-based Key-Value Stores using Shielded Execution.” In: *17th USENIX Conference on File and Storage Technologies (FAST 19)*. Boston, MA: USENIX Association, 2019, pp. 173–190. ISBN: 978-1-931971-48-5. URL: <https://www.usenix.org/conference/fast19/presentation/bailleu>.
- [23] S. Banescu. *Cache Timing Attacks*. July 2011. URL: https://www.academia.edu/3224323/Cache_Timing_Attacks.
- [24] *Benefits of Hardware Trusted Modules*. Accessed: 2019-11-30. URL: <https://www.hardware-security-module.com/benefits/>.
- [25] S. Christopherson. *Intel SGX Virtualization - KVM Forum 2018*. Intel Open Source Technology Center | 01.org Accessed: 2020-10-16. URL: https://www.linux-kvm.org/images/e/e8/KVM_Forum_2018_-_Intel_SGX.pdf.
- [26] V. Costan and S. Devadas. *Intel SGX Explained*. Cryptology ePrint Archive, Report 2016/086. ‘<https://eprint.iacr.org/2016/086>’. 2016.

- [27] V. Costan, I. Lebedev, and S. Devadas. *Sanctum: Minimal Hardware Extensions for Strong Software Isolation*. Cryptology ePrint Archive, Report 2015/564. ' <https://eprint.iacr.org/2015/564/>. 2015.
- [28] *Data leaks: The most common sources*. Accessed: 2020-02-18. URL: <https://www.zdnet.com/pictures/data-leaks-the-most-common-sources/13/>.
- [29] *Data leaks: The most common sources*. Accessed: 2020-02-18. URL: <https://www.zdnet.com/pictures/data-leaks-the-most-common-sources/12/>.
- [30] *Data leaks: The most common sources*. Accessed: 2020-02-18. URL: <https://www.zdnet.com/pictures/data-leaks-the-most-common-sources/14/>.
- [31] G. DeCandia, D. Hastorun, M. Jampani, G. Kakulapati, A. Lakshman, A. Pilchin, S. Sivasubramanian, P. Voshall, and W. Vogels. "Dynamo: Amazon's Highly Available Key-value Store." In: *Proceedings of Twenty-first ACM SIGOPS Symposium on Operating Systems Principles*. SOSP '07. Stevenson, Washington, USA: ACM, 2007, pp. 205–220. ISBN: 978-1-59593-591-5. DOI: 10.1145/1294261.1294281. URL: <http://doi.acm.org/10.1145/1294261.1294281>.
- [32] *Docker*. Accessed: 2020-02-18. URL: <https://www.docker.com>.
- [33] *Docker Compose*. Accessed: 2020-02-18. URL: <https://docs.docker.com/compose/>.
- [34] *Docker Hub Redis Official Repository*. Accessed: 2020-08-15. URL: https://hub.docker.com/_/redis/.
- [35] *Document Store*. Accessed: 2019-06-16. URL: <https://aws.amazon.com/nosql/document/>.
- [36] *FINDING SECRETS...* URL: https://sconedocs.github.io/memory_dump/.
- [37] *Gatling - Load test as code*. URL: <https://gatling.io>.
- [38] C. Gentry, S. Halevi, and N. P. Smart. "Homomorphic Evaluation of the AES Circuit." In: *Lecture Notes in Computer Science*. Springer Berlin Heidelberg, 2012, pp. 850–867. DOI: 10.1007/978-3-642-32009-5_49. URL: https://doi.org/10.1007/978-3-642-32009-5_49.
- [39] *Google Cloud Storage*. Accessed: 2019-07-09. URL: <https://cloud.google.com/storage/>.
- [40] J. Götzfried, M. Eckert, S. Schinzel, and T. Müller. "Cache Attacks on Intel 'SGX'." In: *Proceedings of the 10th European Workshop on Systems Security - 'EuroSec' '17*. 'ACM' Press, 2017. DOI: 10.1145/3065913.3065915. URL: <https://doi.org/10.1145/3065913.3065915>.
- [41] *Graph DBMS*. Accessed: 2019-06-16. URL: <https://aws.amazon.com/nosql/graph/>.

- [42] *Graphene Library OS with Intel SGX Support*. Accessed: 2020-02-16. URL: <https://github.com/oscarlab/graphene>.
- [43] *Graphene-SGX Secure Container (GSC): Automatic Protection of Containerized Applications with Intel SGX*. Accessed: 2020-02-18. URL: <https://github.com/rainfld/gsc>.
- [44] S. Gueron. *A Memory Encryption Engine Suitable for General Purpose Processors*. Cryptology ePrint Archive, Report 2016/204. 2016. URL: <https://eprint.iacr.org/2016/204>.
- [45] S. K. Haider, H. Omar, M. Ahmad, C. Jin, and M. van Dijk. *Intel's SGX In-depth Architecture*. URL: https://scl.engr.uconn.edu/courses/ece6095/lectures/sgx_architecture.pdf.
- [46] *Hardware Trusted Modules*. Accessed: 2019-11-30. URL: <https://resources.infosecinstitute.com/tpms-or-hsms-and-their-role-in-full-disk-encryption-fde/>.
- [47] *How fast is Redis?* Accessed: 2020-02-16. URL: <https://redis.io/topics/benchmarks>.
- [48] M. S. Inci, B. Gulmezoglu, G. Irazoqui, T. Eisenbarth, and B. Sunar. *Seriously, get off my cloud! Cross-VM RSA Key Recovery in a Public Cloud*. Cryptology ePrint Archive, Report 2015/898. 2015. URL: <https://eprint.iacr.org/2015/898>.
- [49] *Intel SGX*. Accessed: 2020-01-04. URL: <https://software.intel.com/en-us/sgx>.
- [50] *intel/kvm-sgx*. Accessed: 2020-10-16. URL: <https://github.com/intel/kvm-sgx>.
- [51] *Intel® Software Guard Extensions (Intel® SGX) - Developer Guide*. Accessed: 2020-01-25. URL: https://download.01.org/intel-sgx/linux-2.2/docs/Intel_SGX_Developer_Guide.pdf.
- [52] *Intel® Software Guard Extensions SDK for Linux*. Accessed: 2020-02-18. URL: <https://01.org/intel-softwareguard-extensions>.
- [53] S. IT. *Key-Value Stores*. Accessed: 2019-06-16. URL: <https://db-engines.com/en/article/Key-value+Stores>.
- [54] S. IT. *Key-Value Stores Ranking*. Accessed: 2019-06-16. URL: <https://db-engines.com/en/ranking/key-value+store>.
- [55] *Jedis - A blazingly small and sane redis java client*. URL: <https://github.com/redis/jedis>.
- [56] *Kernel Virtual Machine*. Accessed: 2020-09-20. URL: https://www.linux-kvm.org/page/Main_Page.

- [57] *Keycloak - Open Source Identity and Access Management*. URL: <https://www.keycloak.org>.
- [58] A. K. Khan and H. J. Mahanta. "Side channel attacks and their mitigation techniques." In: *2014 First International Conference on Automation, Control, Energy and Systems ('ACES')*. IEEE, Feb. 2014. DOI: [10.1109/aces.2014.6807983](https://doi.org/10.1109/aces.2014.6807983). URL: <https://doi.org/10.1109/aces.2014.6807983>.
- [59] T. Kim, J. Park, J. Woo, S. Jeon, and J. Huh. "ShieldStore: Shielded In-memory Key-value Storage with SGX." In: *Proceedings of the Fourteenth EuroSys Conference 2019*. EuroSys '19. Dresden, Germany: ACM, 2019, 14:1–14:15. ISBN: 978-1-4503-6281-8. DOI: [10.1145/3302424.3303951](https://doi.org/10.1145/3302424.3303951). URL: <http://doi.acm.org/10.1145/3302424.3303951>.
- [60] *Kinetic Object Storage*. Accessed: 2019-07-09. URL: <https://storageioblog.com/seagate-kinetic-cloud-object-storage-io-platform/>.
- [61] R. Krahn, B. Trach, A. Vahldiek-Oberwagner, T. Knauth, P. Bhatotia, and C. Fetzer. "Pesos." In: *Proceedings of the Thirteenth EuroSys Conference on - EuroSys 18*. ACM Press, 2018. DOI: [10.1145/3190508.3190518](https://doi.org/10.1145/3190508.3190518). URL: <https://doi.org/10.1145/3190508.3190518>.
- [62] *linux-sgx-driver*. Accessed: 2020-02-18. URL: <https://github.com/intel/linux-sgx-driver>.
- [63] N. Lord. *What is the Principle of Least Privilege (POLP)? A Best Practice for Information Security and Compliance*. Accessed: 2020-10-24. URL: <https://digitalguardian.com/blog/what-principle-least-privilege-polp-best-practice-information-security-and-compliance>.
- [64] F. McKeen, I. Alexandrovich, A. Berenzon, C. Rozas, H. Shafi, V. Shanbhogue, and U. Savagaonkar. *Innovative Instructions and Software Model for Isolated Execution*. Tech. rep. Intel Corporation, Aug. 2013. URL: <https://software.intel.com/sites/default/files/article/413936/hasp-2013-innovative-instructions-and-software-model-for-isolated-execution.pdf>.
- [65] *Memcached*. Accessed: 2019-06-16. URL: <http://www.memcached.org>.
- [66] *Memcached Github*. Accessed: 2019-06-16. URL: <https://github.com/memcached/memcached/wiki/Overview>.
- [67] *Microsoft Azure Cache For Redis*. Accessed: 2019-06-16. URL: <https://azure.microsoft.com/en-us/services/cache/>.
- [68] *Microsoft Azure Cosmos DB*. Accessed: 2019-06-16. URL: <https://azure.microsoft.com/en-us/services/cosmos-db/>.
- [69] *Microsoft Azure Services*. Accessed: 2019-06-16. URL: <https://azure.microsoft.com/en-us>.

BIBLIOGRAPHY

- [70] *Microsoft Security Intelligence Report Volume 22*. Accessed: 2020-01-11. URL: <https://www.microsoft.com/security/blog/2017/08/17/microsoft-security-intelligence-report-volume-22-is-now-available/>.
- [71] M. Minkin. "Improving Performance and Security of Intel SGX." Master's thesis. Israel Institute of Technology, Dec. 2018.
- [72] *Open Redis Servers Infected with Malware*. Accessed: 2020-02-18. URL: <https://www.infosecurity-magazine.com/news/open-redis-servers-infected-with/>.
- [73] *OVHcloud - Intel Software Guard Extensions (SGX)*. Accessed: 2020-02-16. URL: https://www.ovh.ie/dedicated_servers/software-guard-extensions/.
- [74] C. Priebe, K. Vaswani, and M. Costa. "EnclaveDB – A Secure Database using SGX." In: *To appear in the Proceedings of the IEEE Symposium on Security & Privacy, May 2018*. IEEE, May 2018. URL: <https://www.microsoft.com/en-us/research/publication/enclavedb-a-secure-database-using-sgx/>.
- [75] *Professor John McCarthy*. Accessed: 2020-01-11. URL: <https://cs.stanford.edu/memorial/professor-john-mccarthy>.
- [76] *Redis*. Accessed: 2019-06-16. URL: <https://redis.io>.
- [77] *Redis 6.0.06*. Accessed: 2020-08-15. URL: <https://raw.githubusercontent.com/redis/redis/6.0/00-RELEASENOTES>.
- [78] *Redis Multi Model Store*. Accessed: 2019-06-16. URL: <https://db-engines.com/en/system/Redis>.
- [79] *Redis Performance Benchmark*. Accessed: 2019-06-16. URL: <https://redislabs.com/docs/nosql-performance-benchmark/>.
- [80] *Redis Security*. Accessed: 2019-06-16. URL: <https://redis.io/topics/security>.
- [81] *SASL Rfc*. Accessed: 2019-06-16. URL: <https://tools.ietf.org/html/rfc2222>.
- [82] M. Schwarz, S. Weiser, D. Gruss, C. Maurice, and S. Mangard. "Malware Guard Extension: Using SGX to Conceal Cache Attacks." In: *arXiv e-prints*, arXiv:1702.08719 (Feb. 2017), arXiv:1702.08719. arXiv: 1702.08719 [cs.CR].
- [83] *Search Engine Database*. Accessed: 2019-06-16. URL: <https://aws.amazon.com/nosql/search/>.
- [84] *secure-redis-client*. URL: <https://github.com/aanciaes/secure-redis-client>.
- [85] *secure-redis-container*. URL: <https://github.com/aanciaes/secure-redis-container>.
- [86] *secure-redis-proxy*. URL: <https://github.com/aanciaes/secure-redis-proxy>.

-
- [87] *Sgx Memory Limits*. Accessed: 2019-06-16. URL: <https://software.intel.com/en-us/forums/intel-software-guard-extensions-intel-sgx/topic/670322>.
- [88] *SGX Secure Enclaves in Practice: Security and Crypto Review*. Black Hat USA 2016 Security Conference. URL: <https://www.youtube.com/watch?v=0ZVFy4Qsryc>.
- [89] *SGX Virtualization*. Accessed: 2020-10-16. URL: <https://01.org/intel-software-guard-extensions/sgx-virtualization>.
- [90] *SJHOMOLIB*. Accessed: 2020-02-18. URL: <http://nova-lincs.di.fct.unl.pt/prototype/233>.
- [91] M. Taassori, A. Shafiee, and R. Balasubramonian. "VAULT: Reducing Paging Overheads in SGX with Efficient Integrity Verification Structures." In: *ASPLOS '18*. 2018.
- [92] *Tesla cloud systems exploited by hackers to mine cryptocurrency*. Accessed: 2020-02-18. URL: <https://www.zdnet.com/article/tesla-systems-used-by-hackers-to-mine-cryptocurrency/>.
- [93] *The Benefits Of Moving To The Cloud*. Accessed: 2020-01-11. URL: <https://www.forbes.com/sites/forbestechcouncil/2017/05/19/the-benefits-of-moving-to-the-cloud>.
- [94] *The Cambridge Analytica scandal changed the world – but it didn't change Facebook*. Accessed: 2020-02-18. URL: <https://www.theguardian.com/technology/2019/mar/17/the-cambridge-analytica-scandal-changed-the-world-but-it-didnt-change-facebook>.
- [95] *Time Series Databases*. Accessed: 2019-06-16. URL: <https://www.forbes.com/sites/metabrown/2018/03/31/get-the-basics-on-nosql-databases-time-series-databases/>.
- [96] *Trusted Platform Modules*. Accessed: 2019-07-09. URL: <https://docs.microsoft.com/en-us/windows/security/information-protection/tpm/trusted-platform-module-overview>.
- [97] C.-c. Tsai, D. Porter, and M. Vij. "Graphene-SGX: A Practical Library OS for Unmodified Applications on SGX." In: *2017 USENIX Annual Technical Conference (USENIX ATC 17)*. URL: <https://www.usenix.org/system/files/conference/atc17/atc17-tsai.pdf>.
- [98] C.-C. Tsai, D. E. Porter, K. S. Arora, N. Bandi, B. Jain, W. Jannen, J. John, H. A. Kalodner, V. Kulkarni, and D. Oliveira. "Cooperation and security isolation of library OSeS for multi-process applications." In: *Proceedings of the Ninth European Conference on Computer Systems - EuroSys '14*. ACM Press, 2014. DOI: 10.1145/2592798.2592812. URL: <https://doi.org/10.1145/2592798.2592812>.

BIBLIOGRAPHY

- [99] A. Waterman, Y. Lee, D. A. Patterson, K. Asanovic, V. I. U. level Isa, A. Waterman, Y. Lee, and D. Patterson. *The RISC-V Instruction Set Manual*. 2014.
- [100] *What is a Trusted Execution Environment (TEE)?* Accessed: 2019-12-03. URL: <https://www.trustonic.com/news/technology/what-is-a-trusted-execution-environment-tee/>.

TECHNOLOGIES AND VERSIONS

Table I.1 lists and summarises the version of all technologies used in this thesis.

Table I.1: Versions of Used Technologies

Technology	Version
CloudLinux (OVH Cloud OS)	Ubuntu 18.04 LTS Server 64 bits
Linux Kernel	4.15.0-101-generic
Intel SGX Linux	Linux 2.11 Open Source Gold Release
Intel SGX Linux Driver	2.11
Docker	19.03.6
Docker Compose	1.17.1
Redis	6.0rc-1
SCONE	4.2.1
Java OpenJDK 64-Bit	11.0.9
Kotlin	1.4.10
Jedis	3.3.0
HLib	1.2r2
C++	C++14
Linux Musl g++ (GCC)	10.2.0
Spring Boot Starter Web	2.3.0.RELEASE
OpenSSL	1.1.1h
OpenAPI	3.0.0
Keycloak	10.0.2
Scala	2.12.3
Gatling	3.2.1
Jmeter	5.3
Secure-Redis-Container (this)	1
Secure-Redis-Proxy (this)	1
Secure-Redis-Client (this)	1